

ICARUS2 Query Language Specification

Markus Gärtner

2020

Contents

1	JSON-LD Elements	8
1.1	Binding	9
1.2	Constraint	9
1.2.1	Predicate	10
1.2.2	Term	10
1.3	Corpus	10
1.4	Data	11
1.5	Element	11
1.5.1	Sequence	12
1.5.2	Grouping	12
1.5.3	Node	13
1.5.4	Tree Node	13
1.5.5	Edge	13
1.5.6	Element Disjunction	14
1.6	Expression	14
1.7	Group	15
1.8	Import	15
1.9	Lane	16
1.10	Layer	16
1.11	Marker	17
1.11.1	Marker Call	17
1.11.2	Marker Expression	18
1.12	Payload	18
1.13	Property	19
1.13.1	Switches	19
1.14	Quantifier	20
1.15	Query	21
1.16	Reference	21
1.17	Result	22
1.18	Result Instruction	22
1.19	Scope	23
1.20	Sorting	23
1.21	Stream	23
2	Inner ICARUS2 Query Language (IQL) Elements	25
2.1	Reserved Words	25
2.2	Comments	25

2.3	Literals	26
2.3.1	String Literals	26
2.3.2	Boolean Literals	26
2.3.3	Integer Literals	26
2.3.4	Floating Point Literals	27
2.4	Identifiers	27
2.5	Variables and References	28
2.6	Expressions	29
2.6.1	Primary Expressions	29
2.6.2	Path Expressions	30
2.6.3	Method Invocation	30
2.6.4	List Access	30
2.6.5	Annotation Access	31
2.6.6	Type Cast	32
2.6.7	Wrapping	32
2.6.8	Unary Operation	32
2.6.9	Binary Operation	33
2.6.10	Assignment	34
2.6.11	Set Predicate	35
2.6.12	Ternary Operation	36
2.6.13	Value Expansion	36
2.7	Constraints	36
2.8	Payload Structure	37
2.8.1	Filter Constraints	37
2.8.2	Match Modifiers	38
2.8.3	Bindings	38
2.8.4	Selection Statement	39
2.8.5	Lanes	39
2.8.6	Flat Constraints	40
2.8.7	Structural Constraints	40
2.8.8	Sequence Constraints	47
2.8.9	Tree Constraints	48
2.8.10	Graph Constraints	49
2.8.11	Global Constraints	49
2.9	Result Processing	50
3	Utility Markers & Functions	51
3.1	Position Markers	51
3.1.1	Negation	51
3.1.2	Reverse Indices	52
3.1.3	Relative Indices	52
3.1.4	Marker Stacking	53
3.1.5	Markers & Quantification	53
3.1.6	Markers in Global Constraints	54
3.2	Tree Markers	54
3.2.1	Hierarchy Markers	54
3.2.2	Path Markers	57

3.2.3	Evaluation Performance	58
3.3	Sequence Functions	58
3.4	Spatial Functions	59
3.5	Tree Functions	60
3.6	Graph Functions	60
Appendices		61
A	Extended Grammar Diagrams	62

List of Figures

2.1	Illustration of quantification	44
2.2	Example target tree for simple query	49
3.1	Fringe illustration	57

List of Tables

1.1	Supported query switches	20
2.1	Identifier types	28
2.2	Binary operators	33
2.3	Boolean conversion rules	37
2.4	Node quantifiers	44
3.1	Global position markers	52
3.2	Tree positional markers	55
3.3	Hierarchical markers	56
3.4	Path markers	57
3.5	Sequence functions	59
3.6	Spatial functions	59
3.7	Tree functions	60

Todo list

say something about the namespace and general iql: prefixing	8
expand on the order of matches	49
content, explain node and edge composition, mention referencing as a strategy to mini- mize edge declaration overhead, etc...	49
list properties required for hoisting and mention hoisting of constraint fragments, since IQL splits boolean expressions into fragments	50
rework, since we dropped the DNF requirement	53
finish	58
correct description/caption in table	58

Introduction

Queries in the IQL are designed to be self-contained with logical sections for specifying all the information required to determine the target of a query and its granularity, resolve additional dependencies such as extensions or scripts, link and validate constraints to parts of the target corpus or corpora and finally optional pre- and post-processing steps. To achieve this complex task IQL embeds a keyword-based syntax for the query payload within a JSON-LD structure to drive declaration of all the aforementioned information. As a side effect queries can become quite verbose and potentially cumbersome to define manually. As a countermeasure the overall structure of a query is composed of blocks that can be glued together incrementally and that make it very easy for an application built on top of it to provision boilerplate query code based on settings or a GUI so that the user only needs to type the actual constraints used in the query (the so called *query payload*, cf. Chapter 2). This document lists the basic building blocks of queries and their compositions.

Chapter 1 gives an overview on the JSON-LD elements available in the “outer” section of the IQL protocol. Chapter 2 provides the specification for the actual query language used in a query’s “inner” payload and Chapter 3 finally contains various collections of utility functions available to make queries both simpler to declare and more efficient to evaluate by the ICARUS2 Corpus Query Processor (ICQP).

Chapter 1

JSON-LD Elements

The “outer layer” of every serialized IQL query is a JSON-LD object with various (optional) properties to hold all the required information for proper query evaluation.

The following snippet shows the mockup of a simple query that searches the TIGER corpus for word pairs starting with an adjective and ending in a word with the surface form “test”, ignoring case and only returning up to 100 hits in a keyword-in-context (KWIC) style. The query also features an import statement for the STTS part-of-speech tagset, allowing more controlled expressions inside query constraints.

```
1 { "@context" : "http://www.ims.uni-stuttgart.de/icarus/v2/jsonld/iql/  
  query"  
2   "@type" : "iql:Query",  
3   "@id" : "query/000001",  
4   "iql:imports" : [ {  
5     "@type" : "iql:Import",  
6     "@id" : "import/000001",  
7     "iql:name" : "common.tagsets.stts"  
8   } ],  
9   "iql:setup" : [ {  
10    "@type" : "iql:Property",  
11    "iql:key" : "iql.string.case.off"  
12  } ],  
13  "iql:streams" : [ {  
14    "@type" : "iql:Stream",  
15    "@id" : "stream/000001",  
16    "iql:corpus" : {  
17      "@type" : "iql:Corpus",  
18      "@id" : "corpus/000001",  
19      "iql:name" : "TIGER-v2"  
20    },  
21    "iql:rawPayload" : "FIND ADJACENT [pos==stts.ADJ][form==\"test\"]",  
22    "iql:result" : {  
23      "@type" : "iql:Result",  
24      "iql:resultTypes" : [ "kwic" ],  
25      "iql:limit" : 100  
26    }  
27  } ]  
28 }
```

1.1 Binding

A binding associates a collection of member variables (2.5) with the content of a specific item layer or derived layer type.

Attributes of `iql:Binding`:

Attribute	Type	Required	Default
distinct	Boolean	no	false
edges	Boolean	no	false
target	string	yes	

`iql:distinct` Enforces that the bound member references in this binding do **not** match the same target items during evaluation. Depending on the structural constraints used in the query, this setting might be redundant (e.g. when using the member references as identifiers for tree nodes who already are structurally distinct), but can still be used to make that fact explicit.

`iql:edges` Signals that the member labels are to be used for edges within a structure.

`iql:target` The name or alias of the layer to whose content the member variables should be bound.

Nested Elements of `iql:Binding`:

Element	Type	Required
members	array of <code>iql:Reference</code> (1.16)	yes

`iql:members` Non-empty collection of member references that are bound to the target layer's content. Every such instance of `iql:Reference` (1.16) must be unique within the surrounding `iql:Payload` (1.12).

1.2 Constraint

Constraints represent the actual content filtering of every query.

Attributes of `<Constraint>`:

Attribute	Type	Required	Default
id	string	yes	
solved	Boolean	no	false
solvedAs	Boolean	no	false

`iql:id` Identifier to uniquely identify the constraint within the entire query.

`iql:solved` Hint for the evaluation engine that this constraint has already been solved, either by a back-end implementation or as a result of (partial) query evaluation by the engine itself.

`iql:solvedAs` Specifies to what Boolean value (**true** or **false**) the constraint has been evaluated.

1.2.1 Predicate

Wraps a Boolean **iql:Expression** into an atomic constraint element that represents the smallest unit of evaluation for the top-level evaluation engine.

Extends <Constraint>(1.2).

Nested Elements of iql:Predicate:

Element	Type	Required
expression	iql:Expression (1.6)	yes

iql:expression The actual expression to be evaluated to a Boolean result. Note that typically this expression **cannot** be composed of directly nested Boolean conjunctions or disjunctions, as the engine will have parsed those into **iql:Term** (1.2.2) objects already during the first processing phase.

1.2.2 Term

A collection of constraints with a logical connective.

Extends <Constraint>(1.2).

Attributes of iql:Term:

Attribute	Type	Required	Default
operation	enum	yes	

iql:operation The Boolean connective to be applied to all the constraint items. Legal values are the strings “conjunction” or “disjunction”.

Nested Elements of iql:Term:

Element	Type	Required
items	array of iql:Constraint (1.2)	yes

iql:items The constraints which are to be combined by the specified **iql:operation**.

1.3 Corpus

Top-level entry point for querying a single stream.

Attributes of iql:Corpus:

Attribute	Type	Required	Default
id	string	yes	
name	string	yes	
pid	string	no	

iql:id Identifier to uniquely identify the corpus within the entire query.

iql:name The identifier used by the query engine’s manifest registry for the corpus.

iql:pid Optional external identifier used for the corpus. This can be any persistent identifier such as a DOI, VLO-Handle or similar means of referencing a resource.

1.4 Data

Allows to embed binary data in the query and make it usable from within constraint expressions via a designated reference.

Attributes of **iql:Data**:

Attribute	Type	Required	Default
id	string	yes	
name	string	yes	
content	string	yes	
codec	string	no	hex
checksum	string	no	
checksumType	enum	no	

iql:id Identifier to uniquely identify the corpus within the entire query.

iql:name The identifier used for the expression (2.6) which can be used to reference the binary payload from within query constraints.

iql:content The actual content of the payload in textual form. How to properly convert the textual form to a binary stream is defined by the **iql:codec** attribute.

iql:codec Specifies the mechanism of converting the **iql:content** data into an actual binary stream. If left empty, defaults to **hex**.

iql:checksum Optional hex-string of the checksum to check the **iql:content** against.

iql:checksumType Defines the algorithm for computing the checksum. Currently only **MD5** is supported as legal value.

1.5 Element

Abstract base type for all logical and/or structural units that can be matched against content of a target corpus.

Attributes of **<Element>**:

Attribute	Type	Required	Default
id	string	yes	
consumed	Boolean	no	false

iql:id Identifier to uniquely identify the element within the entire query.

iql:consumed Signals that the element has already been *used up* in the context of a partial query evaluation. An element that has been consumed can safely be ignored in the further evaluation of the query. Note that this state can be propagated according to the following rules:

- An **iql:Node**(1.5.3) can be marked as **consumed** if its **iql:constraint** is marked as **solved** and its match count satisfies the **iql:quantifiers** requirement. Note that cross-referencing constraints can only be considered solved when all other aspects of the involved elements support the **consumed** state.

- An **iql:TreeNode**(1.5.4) can be marked as **consumed** if above conditions are met and all nested **iql:children** are marked **consumed**.
- An **iql:Edge**(1.5.5) is considered **consumed** when both its terminals are **consumed** and the same conditions regarding its **iql:constraint** are fulfilled as mentioned above.
- An **iql:ElementDisjunction**(1.5.6) is considered **consumed** if at least one of its **iql:alternatives** has been marked **consumed**.

1.5.1 Sequence

Wrapper around a list of **iql:Element** (1.5) instances to group them for either nesting or disjunction.

Extends iql:Element(1.5).

Attributes of iql:NodeSet:

Attribute	Type	Required	Default
arrangement	enum	no	unspecified

iql:arrangement Defines what kind of order or arrangement should be assumed between the elements in this sequence. Legal values are **unspecified**, **ordered** (matched elements must occur in exactly the order specified in this set but need not form a continuous span) or **adjacent** (matched elements must form a continuous span).

Nested Elements of iql:NodeSet:

Element	Type	Required
nodes	array of <Element> (1.5)	no

iql:nodes List of nested <Element> instances. Legal types depend on the context in which this sequence is being used.

1.5.2 Grouping

Wrapper around an **iql:Element** (1.5) instance to allow for additional grouping and/or quantification.

Extends iql:Element(1.5.3).

Nested Elements of iql:ElementGrouping:

Element	Type	Required
element	<Element> (1.5)	no
quantifiers	array of iql:Quantifier (1.14)	no

iql:element Nested <Element> instance. Legal types depend on the context in which this grouping is being used.

iql:quantifiers Optional quantifiers to define the multiplicity of matches of this grouping required for a positive evaluation. Multiple quantifiers behave disjunctively.

1.5.3 Node

Logical unit for sequence or graph matching in a target corpus. May contain local constraints and can also be quantified.

Extends <Element>(1.5).

Attributes of iql:Node:

Attribute	Type	Required	Default
label	string	no	

iql:label Identifier to bind the node through a previously defined **iql:Binding** (1.1) declaration.

Nested Elements of iql:Node:

Element	Type	Required
constraint	<Constraint> (1.2)	no
quantifiers	array of iql:Quantifier (1.14)	no

iql:constraint Optional local constraint to be matched against the content of potential target candidates during query evaluation.

iql:quantifiers Optional quantifiers to define the multiplicity of matches of this node required for a positive evaluation. Multiple quantifiers behave disjunctively.

1.5.4 Tree Node

Extension of the simple **iql:Node** type (1.5.3) to add implicit hierarchical constraints related to dominance within tree structures.

Extends **iql:Node**(1.5.3).

Nested Elements of iql:TreeNode:

Element	Type	Required
children	instance of <Element> (1.5)	no

iql:children Optional nested nodes or node alternatives.

1.5.5 Edge

Specialized element extension to query structural information in graphs.

Extends <Element>(1.5).

Attributes of iql:Edge:

Attribute	Type	Required	Default
label	string	no	
edgeType	enum	yes	

iql:label Identifier to bind the edge through a previously defined **iql:Binding** (1.1) declaration.

iql:edgeType The type specification for this edge, primarily a directionality information. Legal values are **simple**, **one-way** or **two-way**.

Nested Elements of `iql:Edge`:

Element	Type	Required
constraint	<Constraint> (1.2)	no
source	<code>iql:Node</code> (1.5.3)	yes
target	<code>iql:Node</code> (1.5.3)	yes

`iql:constraint` Optional local constraint to be matched against the content of potential target candidates during query evaluation.

`iql:source` Source node declaration.

`iql:target` Target node declaration.

For complex graph declarations multiple nodes can be defined having the same `iql:label`. The evaluation engine will treat them as being the same node. Note however, that at most **one** node per label is allowed to declare a local `iql:constraint` attribute!

1.5.6 Element Disjunction

Allows declaration of multiple alternative element definitions. When evaluating the query, each such alternative that is matched successfully will cause this element declaration to evaluate positively.

Extends <Element>(1.5).

Nested Elements of `iql:ElementDisjunction`:

Element	Type	Required
alternatives	array of <code>iql:Element</code> (1.5)	yes

`iql:alternatives` The alternative element declarations, each of which constitutes a legal match for this element declaration. Must not contain less than 2 elements!

1.6 Expression

Wraps the textual form of an arbitrarily complex IQL expression, which can be a formula, literal, method invocation, a combination of those or a great many other types of expressions. For more details see Section 2.6.

Attributes of `iql:Expression`:

Attribute	Type	Required	Default
content	string	yes	
resultType	string	no	

`iql:content` The textual form of the expression. Must be valid according to the specifications in Section 2.6.

`iql:resultType` An optional specification regarding the return type of the expression. Redundant when the expression is used as a constraint, as those are required to always evaluate to a Boolean result value anyway.

1.7 Group

Provides a mechanism to collect successful matches into dedicated groups, either for result visualization or use in further result processing.

Attributes of `iql:Group`:

Attribute	Type	Required	Default
id	string	yes	
label	string	yes	

`iql:id` Identifier to uniquely identify the group declaration within the entire query.

`iql:label` Label (ideally human readable) to be used for referencing this group in subsequent result processing or for generating textual result reports.

Nested Elements of `iql:Group`:

Element	Type	Required
groupBy	<code>iql:Expression</code> (1.6)	yes
filterOn	<code>iql:Expression</code> (1.6)	no
defaultValue	<code>iql:Expression</code> (1.6)	no

`iql:groupBy` The mandatory expression used to extract the value from matches based on which the actual grouping occurs.

`iql:filterOn` Optional mechanism to exclude certain matches from being used for grouping.

`iql:defaultValue` If matches cannot produce a valid value for grouping but should still be included in the process, this optional field provides the means of declaring a kind of “fall back” group. Be aware of potential overlap in groups when using default values that are not distinct from the regular grouping results.

1.8 Import

To allow for flexible integration of macro definitions or bigger language extensions, IQL provides an optional section in the query that lets users specify exactly what additional modules besides the bare IQL core are required for evaluating the query. Each import target is specified by providing it's unique name and telling the engine whether or not the import is to be considered optional.

Attributes of `iql:Import`:

Attribute	Type	Required	Default
id	string	yes	
name	string	yes	
optional	Boolean	no	false

`iql:id` Identifier to uniquely identify the import within the entire query.

`iql:name` The original name of the extension to be added.

`iql:optional` Defines whether or not the referenced extension is optional. Non-optional imports that cannot be resolved to an actual extension during the query evaluation phase will cause the entire process to fail.

1.9 Lane

Lanes serve as a means of splitting queries for a single corpus stream into multiple logical subqueries that target different structural and/or logical layers, e.g. multiple syntactic analyses for the same source text.

Attributes of `iql:Lane`:

Attribute	Type	Required	Default
id	string	yes	
name	string	yes	
alias	string	no	
laneType	enum	yes	

iql:id Identifier to uniquely identify the lane within the entire query.

iql:name The unique identifier of the item layer or structure layer that serves as target for this lane.

iql:alias If items of this lane in their entirety are meant to be used as part of query expressions inside this field holds the label used for the respective member variable. It is recommended to keep the chosen **alias** close to the original **name** to avoid confusion.

iql:laneType The type of structure this lane is meant to match, effectively defining the basic complexity class for evaluation. legal values are **sequence**, **tree** and **graph**. Note that the initial evaluation engine for IQL does not support the **graph** type!

Nested Elements of `iql:Lane`:

Element	Type	Required
elements	<Element> (1.5)	yes

iql:elements The structural constraints to be used for evaluation of this lane.

1.10 Layer

Every layer selector either references an entire subgraph of the corpus' member-graph directly or constructs a partial selection as part of a **iql:Scope** (1.19). When using the first approach, an item layer is referenced and all its dependencies and associated annotation layers will be made available implicitly. This is an easy way of accessing simple corpora, but can lead to costly I/O overhead when loading vast parts of a complex corpus that aren't actually needed to evaluate the query. For a more fine-grained alternative, scopes allow to create a scope that spans an exactly specified collection of layers. If multiple layer selectors are defined, up to one can be declared as "primary" to represent the granularity of returned items for the search or scope. In case no layer is explicitly marked as "primary", the one specified by the corpus or context will be used for that role by default.

Attributes of `iql:Layer`:

Attribute	Type	Required	Default
id	string	yes	
name	string	yes	
alias	string	no	
primary	Boolean	no	false
allMembers	Boolean	no	false

iql:id Identifier to uniquely identify the layer within the entire query.

iql:name Identifier used to reference the layer within its host corpus.

iql:alias Optional identifier to rename the layer for referencing within the query.

iql:primary Signals that the layer is intended to act as the primary layer in the query or scope and as such defines the level of granularity for obtaining chunks in the corpus.

iql:allMembers When this layer definition is used inside a **iql:Scope** (1.19), effectively adds the entire member-subgraph of this layer to the scope. This property is redundant when the layer is part of the regular **iql:layers** declaration in a **iql:Stream** (1.21), as in that case all member subgraphs for each layer are already being added to the global scope!.

1.11 Marker

Abstract base type for all positional, hierarchical or other sorts of (tree) markers available in IQL.

1.11.1 Marker Call

Identifies a marker by name and optionally provides a set of numerical arguments for instantiating the marker.

Extends <Marker>(1.11).

Attributes of iql:MarkerCall:

Attribute	Type	Required	Default
name	string	yes	

iql:name The name to uniquely identify the marker. Although in their introduction chapter 3 they appear camel-cased, actual resolution of markers by name is done case-insensitive.

Nested Elements of iql:Payload:

Element	Type	Required
arguments	array of numbers	no

iql:arguments Optional array of numerical arguments for the marker. Note that markers support both absolute (integer) and relative (floating point) index arguments, as well as negative ones to express reverse indices.

1.11.2 Marker Expression

Used to group multiple markers into boolean expressions.

Extends <Marker>(1.11).

Attributes of iql:MarkerExpression:

Attribute	Type	Required	Default
expressionType	enum	yes	

iql:expressionType Defines whether the group should be treated as a **disjunctive** or **conjunctive** expression. **disjunctive**

Nested Elements of iql:Payload:

Element	Type	Required
items	array of <Marker>(1.11)	yes

iql:items List of the nested markers to be evaluated.

1.12 Payload

Every payload encapsulates all the (processed) query constraints to be evaluated against a single stream of corpus data.

Attributes of iql:Payload:

Attribute	Type	Required	Default
id	string	yes	
name	string	no	
queryType	enum	yes	
queryModifier	enum	no	

iql:id Automatically generated identifier to uniquely identify the payload within the entire query.

iql:name Custom identifier to uniquely identify the payload within the entire query. This attribute is deprecated but currently being kept to shift its use case.

iql:queryType The overall type of query strategy to be applied for this query payload. Legal values are **all** (returns the entire corpus and disallows any kind of constraint, leaving only the **iql:Result** (1.17) declaration as option to modify the result volume), **plain** (disabling any kind of structural constraints/lanes), **singleLane** and **multiLane**. The last two values dictate the minimal/maximal number of **iql:Lane** definitions in this payload.

iql:queryModifier Allows to limit the number of times an individual units-of-interest (UoIs) will be returned in the result. Supported values are **first**, **last** and **any**. The specific semantics of this modifier are described in more details in Section 2.7.

Nested Elements of **iql:Payload**:

Element	Type	Required
bindings	array of iql:Binding (1.1)	no
lanes	array of iql:Lane (1.9)	no
filter	iql:Constraint (1.2)	no
constraint	iql:Constraint (1.2)	no

iql:bindings Optional collection of bindings used within this payload. Note that member variables inside constraints or structural query elements will not resolve unless previously bound to corpus members.

iql:lanes If **iql:queryType** is set to **singleLane** or **multiLane**, this array is expected to hold either exactly 1 or at least 2 **iql:Lane** declarations that define structural constraint for the evaluation.

iql:filter If **iql:queryType** is set to anything other than **plain**, this constraint expression allows to filter contextual UoIs prior to the actual structural matching.

iql:constraint If **iql:queryType** is set to **plain**, this is expected to contain the basic constraints for matching candidates. In any version involving **iql:Lane** declarations, global constraints can be defined here as a means of implementing complex query features that are tested once the lanes have produced preliminary result candidates.

1.13 Property

Allows customization of the evaluation process by changing parameters or switching certain features on/off.

Attributes of **iql:Property**:

Attribute	Type	Required	Default
key	string	yes	
value	string	no	

iql:key The identifier of the targeted parameter or switch. The evaluation engine might report unknown keys as errors.

iql:value The actual value to apply to the specified property in case it is not a switch.

1.13.1 Switches

For increased flexibility, IQL supports a collection of switches to turn certain optional features on or off when needed. Switches are static and cannot be changed for the active query evaluation once set. All the native IQL switches use the prefix **iql:** for their name. Any extensions that offer additional switches should declare and use their own namespace for those switches! Currently supported switches are shown in Table 1.1.

Name	Description
iql.string.case.off	Turns of case sensitivity when performing string operations such as equality checks.
iql.string.case.lower	Another approach to case insensitivity, this switch turns all strings into lower case.
iql.expansion.off	Effectively shuts down value expansion Section 2.6.13.
iql.string2bool.off	Deactivates the interpretation of strings as Boolean values as described in Section 2.7.
iql.int2bool.off	Deactivates the interpretation of integers as Boolean values as described in Section 2.7.
iql.float2bool.off	Deactivates the interpretation of floating point numbers as Boolean values as described in Section 2.7.
iql.obj2bool.off	Deactivates the interpretation of arbitrary objects as Boolean values as described in Section 2.7.
iql.any2bool.off	Deactivates the interpretation of anything non-Boolean as Boolean value. This is a combination of “iql.string2bool.off”, “iql.int2bool.off”, “iql.float2bool.off” and “iql.obj2bool.off”.
iql.direction.reverse	Reverses the direction used to traverse top-level corpus data for a search.
iql.array.zero	Change array access (2.6.4) to be 0-based.
iql.markers.position.relative	Allow position markers (3.1) to use relative (percentage) arguments.
iql.warnings.off	Deactivates all warnings, potentially resulting in confusing results if there are mistakes in the query.
iql.parall.off	Forces the query evaluation engine to run single-threaded. This does however only affect the actual matcher, not additional. modules such as monitoring or item caches

Table 1.1: Currently supported switches in IQL and their explanations.

1.14 Quantifier

Specifies the multiplicity of an associated <Element> (1.5).

Attributes of iql:Quantifier:

Attribute	Type	Required	Default
quantifierType	enum	yes	
value	integer	no	
lowerBound	integer	no	
upperBound	integer	no	

iql:quantifierType Defines how to interpret the other attributes. Legal values are **all** (universal quantification), **exact**, **atMost** (0..n), **atLeast** (n+), **range** (n..m).

iql:value Target or limit value when **iql:quantifierType** is set to **exact**, **atMost** or **atLeast**.

iql:lowerBound Used for **range** quantification to define the minimum multiplicity.

iql:upperBound Used for **range** quantification to define the maximum multiplicity.

1.15 Query

Encapsulates all the global configuration and extension of the query engine, as well as shared embedded data. Each query contains at least one **iql:Stream** declaration that in turn holds the actual query payload with constraints for the matching process.

Attributes of iql:Query:

Attribute	Type	Required	Default
id	string	yes	
dialect	string	no	1.0

iql:id Identifier for the query, chosen by the client. In more complex (asynchronous) query workflows this id is used to map answers and results to the correct query.

iql:dialect Specifies which basic version of IQL to use. The initial version of IQL is “1.0” and by leaving the dialect part of a query blank the engine will default to this initial version.

Nested Elements of iql:Query:

Element	Type	Required
imports	array of iql:Import (1.8)	no
setup	array of iql:Property (1.13)	no
embeddeData	array of iql:Data (1.4)	no
streams	array of iql:Stream (1.21)	yes

iql:imports Defines extensions to be applied to the evaluation engine prior to actual query evaluation.

iql:setup Allows to configure the core evaluation engine or already defined extensions in a simple manner.

iql:embeddeData Binary data to be used in the evaluation process, such as audio or video fragments.

iql:streams Corpus data streams to be queried. In the initial version, the engine only supports single-stream querying!

1.16 Reference

Models references usable from within query expressions for accessing corpus members or variables.

Attributes of iql:Reference:

Attribute	Type	Required	Default
id	string	yes	
name	string	yes	
referenceType	enum	yes	

iql:id Identifier to uniquely identify the reference within the entire query.

iql:name The local identifier to be used for addressing this reference. Note that this is the bare name without any type-specific prefixes (such as '\$' for members, cf. Section 2.5).

iql:referenceType Specifies the nature of this reference. Legal values are **reference**, **member** or **variable**.

1.17 Result

Encapsulates all the information on result processing and preparation.

Attributes of iql:Result:

Attribute	Type	Required	Default
limit	integer	no	
percent	Boolean	no	false

iql:limit Optional limitation on the total size of the result to be returned. If the **iql:percent** flag is not set to **true**, this number is in reference to the units provided by the query's primary item layer.

iql:percent If set to **true** the value defined in **iql:limit** is treated as a integer percentage value in the interval 1 to 99, with boundaries included.

Nested Elements of iql:Result:

Element	Type	Required
resultTypes	array of enum	yes
resultInstructions	array of iql:ResultInstruction (1.18)	no
sortings	array of iql:Sorting (1.20)	no

iql:resultTypes Defines the result format or type the engine should return data in. At least one result type must be declared and the engine can also be instructed to return the results in multiple formats simultaneously. In the first iteration only **kwic** (keyword-in-context) and **custom** (as a placeholder for the raw corpus members) are supported.

iql:resultInstructions Optional collection of additional processing instructions to generate (textual) result reports.

iql:sortings Allows to sort matches before generating result reports.

1.18 Result Instruction

Currently unused dummy for declaring post-processing instructions on the query result to perform conversions and/or tabular calculations.

1.19 Scope

Very detailed vertical filtering of the layers available in a query.

Attributes of `iql:Scope`:

Attribute	Type	Required	Default
id	string	yes	

`iql:id` Identifier to uniquely identify the scope within the entire query.

Nested Elements of `iql:Scope`:

Element	Type	Required
layers	array of <code>iql:Layer</code> (1.10)	yes

`iql:layers` The layer members of this scope.

1.20 Sorting

Defines a single rule for sorting query results based on an arbitrarily complex expression.

Attributes of `iql:Sorting`:

Attribute	Type	Required	Default
order	enum	yes	

`iql:order` Hint on sorting direction, legal values are **`asc`** or **`desc`** for ascending or descending order, respectively.

Nested Elements of `iql:Sorting`:

Element	Type	Required
expression	<code>iql:Expression</code> (1.6)	yes

`iql:expression` The actual sorting expression. It can use any (member) reference or variable available in the query to compute its result and must return a type that is comparable to allow stable sorting. Per default any of the primitive numerical types (**`int`** or **`float`**), the text type **`string`** and any member of the ICARUS2 framework implementing the **`java.lang.Comparable`** interface can be used as return type.

1.21 Stream

A stream encapsulates all the information and query constraints to extract, evaluate and prepare data from a single corpus. Note that many of the attributes and/or elements below are marked as optional, but the following restrictions are in effect:

- Either **`iql:rawPayload`** or **`iql:payload`** must be provided by the client.
- Either **`iql:layers`** or **`iql:scope`** must be provided to define the granularity of data being loaded for evaluation.

Attributes of iql:Stream:

Attribute	Type	Required	Default
id	string	yes	
primary	Boolean	no	false
rawPayload	string	no	
rawGrouping	string	no	
rawResult	string	no	

iql:id Identifier to uniquely identify the stream within the entire query.

iql:primary Flag to indicate that the primary layer of this stream is meant to be used as primary layer of the entire search result. Only one stream can declare this property and it primarily dictates the order of result elements in a multi-stream query or which stream is allowed to dictate sorting.

iql:rawPayload The textual (raw) form of the payload for this stream, i.e. all the constraints and structural query content.

iql:rawGrouping The textual (raw) grouping definitions to be applied for results of this stream.

iql:rawResult The textual (raw) result configuration and post-processing instructions for this stream.

Nested Elements of iql:Stream:

Element	Type	Required
corpus	iql:Corpus (1.3)	yes
layers	array of iql:Layer (1.10)	no
scope	iql:Scope (1.19)	no
payload	iql:Payload (1.12)	no
grouping	array of iql:Grouping (1.7)	no
result	iql:Result (1.17)	yes

iql:corpus The corpus to extract data from.

iql:layers Vertical filtering to be applied to the corpus prior to actual query evaluation.

iql:scope Another and more fine-grained form of vertical filtering that allows for more precise selection of layers to be part of this stream's data.

iql:payload The processed form of **iql:rawPayload**.

iql:grouping The processed form of **iql:rawGrouping**.

iql:result The processed form of **iql:rawResult**.

Chapter 2

Inner IQL Elements

Certain parts of an IQL query can be defined in *raw* form, that is, in a keyword-driven formal language. During the first phase of query evaluation they get (partly) translated into their respective JSON-LD counterparts described in Chapter 1 (unless of course the query or query fragments are provided fully processed). This section defines the syntax and additional rules for those raw statements. Note that the textual form of all following IQL elements is expected to be encoded in UTF-8, so no special escape mechanisms are needed for unicode content.

2.1 Reserved Words

The following list of keywords is reserved and any of the words may not be used as direct identifier strings in a query. They are reserved in both all lowercase and all uppercase variants, and while camel-cased versions are technically permitted, it is highly discouraged to use them:

ADJACENT	DISTINCT	GROUP	OMIT
ALL	DO	HAVING	ON
AND	EDGES	HITS	OPTIONAL
ANY	END	IN	OR
AS	EVEN	LABEL	ORDER
ASC	FALSE	LANE	ORDERED
BY	FILTER	LAST	RANGE
COUNT	FIND	LIMIT	STEP
DEFAULT	FIRST	NOT	TRUE
DESC	FOREACH	NULL	UNORDERED
DISJOINT	FROM	ODD	WITH

In addition the following strictly lowercase words are reserved as type identifiers and may not be used otherwise:

boolean	int	float	string
---------	-----	-------	--------

2.2 Comments

IQL supports single-line comments, indicated by “//”. All remaining content in a line after the comment indicator will be ignored when parsing and evaluating a query.

2.3 Literals

Literals are statically-typed fixed-value expressions in IQL. They are parsed only once during the initial processing part of a query.

2.3.1 String Literals

IQL uses simple double quotes ("" or U+0022) to define string literals. String literals may not contain any of the following symbols directly:

```
\n line break
\r carriage return
\t tab
\ backslash
" nested quotation mark
```

Any of those symbols listed above can be embedded into a string literal as part of an escape sequence with a preceding backslash. At the current time there is no planned mechanism to provide additional escape support for unicode symbols, since the default encoding scheme for IQL is UTF-8.

Examples for valid string literals:

```
"string"
"123"
"some fancy number (123.456e-789) and special symbol ♣"
"a more complex string!"
"a\n multiline\n string..."
```

2.3.2 Boolean Literals


Boolean literals are limited to either all lowercase or all uppercase versions of the literals **true** and **false**.

2.3.3 Integer Literals

Signed Integer Literals Literals representing regular **int** (32bit) or **long** (64bit) integers consist of an optional initial sign ('+' or '-') and the body consisting of digits ('0' to '9') or underscore ('_') characters. Underscore characters may only appear inside the integer literal, never at the beginning or end (not counting the sign symbol).

Grammar Snippet 2.1

$\langle integerLiteral \rangle$:



Examples for valid (signed) integer literals:

1
+123
-123
1_000_000
-99_000000_0

Pure Integer Literals Some parts of the IQL syntax only allow unsigned "pure" integers and will explicitly state this fact. In those special cases integer literals may neither contain the initial sign symbol nor intermediate underscores.

2.3.4 Floating Point Literals

Floating point literals are constructed by using a (signed) integer literal for the pre-decimal part, a dot '.' as delimiter and a decimal part made up by a unsigned integer literal. They represent either single-precision **float** (32bit) or double-precision **double** (64bit) values.

Grammar Snippet 2.2

$\langle floatingPointLiteral \rangle:$
 $\rightarrow \langle signedInteger \rangle - '.' - \langle unsignedInteger \rangle \rightarrow$

Examples for valid (signed) floating point literals:

1.0
+123.456
-123.456
1_000_000.999
-99_000000_0.000_000_001

While many languages offer to express floating point literals in the scientific notation with explicit exponent declaration, we do not include this in the initial draft of IQL.

2.4 Identifiers

Identifiers in IQL are combinations of lowercase or uppercase alphabetic [**a-zA-Z**] characters that may contain underscore symbols **_** between the first and last position and may also contain digits [**0-9**] on any position except as initial symbol (cf. snippet 2.3).

Grammar Snippet 2.3

$\langle identifier \rangle:$
 $\rightarrow \langle character \rangle$

The diagram illustrates the recursive definition of an identifier. It starts with a horizontal arrow representing the sequence of characters. Below the arrow, a bracket indicates that the sequence can be expanded into a character followed by another identifier. This process repeats, showing that an identifier is a character followed by zero or more occurrences of a character or digit. The final expansion shows a character followed by a digit, which is then followed by an underscore, indicating that underscores are allowed between characters and digits.

Examples for valid identifiers:

```
x
myIdentifier
x1
x_1
x__1
x321
some_random_id
someRandomId002
random_2_4
notTheBest_____example
```

Identifiers are limited in length by the engine to a total of 255 characters. This is a purely arbitrary choice to keep queries readable and not subject to any technical limitations.

2.5 Variables and References

In IQL all top-level (i.e. not part of the tail expression in a hierarchical path) identifiers are expected to reference 'something' from the global namespace available to the query. This namespace is populated with all the globally available constants, methods and helper objects from the IQL core and any imported extensions, as well as all the corpus members defined in the scoping part of the query. Outside this global namespace any dynamically created identifiers from within a query reside in the variable namespace and are marked with a preceding '@' (e.g. `@myVariable`). They can be used the same way as any regular identifier, with the exception of additionally allowing assignment expressions when inside script blocks. In addition any corpus members bound within a constraint section are prefixed with a '\$' sign, such as `$token1`. Table 2.1 provides a compact overview of the available identifiers and their capabilities/features.

Type	Prefix	Example	Scope	Fixed ¹	Final	Re-Assign
Reference		<code>max()</code>	global	X	X	
Variable	@	<code>@myVar</code>	limited	(X)		X
Member	\$	<code>\$token</code>	limited	X	(X)	

Table 2.1: Identifier types available in IQL and their properties.

Special remarks: Variables are more or less general-purpose storage objects for arbitrary values and without a fixed type. Their first assignment however hints at the implied type to be used and as such they can cause cast errors when used for situations where an incompatible type would be needed.

Member identifiers are final in the sense that they cannot be re-assigned explicitly but will be implicitly for every iteration of the query on a new part of the corpus. For example, above `$token` member will point to a new token object every time the inner constraint parts of the query are evaluated. Therefore member identifiers could be viewed as a sort of loop variable.

2.6 Expressions

Expressions are the foundation of every query. Each expression has a (usually fixed) result type and evaluates to a value of that type. They can take any of the following forms and an overview diagram is available in the appendix as snippet A.1.

2.6.1 Primary Expressions

Any literal of types **boolean**, **string**, **int** or **float** can serve as a primary expression of that type. See Section 2.3 for examples and a more detailed specification of the various types of literals in IQL.

2.6.1.1 References

Any reference or variable as described in Section 2.5 is also a valid primary expression. Note that the process for resolving references strongly relies on the context the reference is being used in, for instance to allow simple references to be used as aliases for method calls.

Additionally a special construct is available as primary expression to more easily access annotation values from within a node definitions (2.8.7). Normally the evaluation engine tries to resolve any otherwise unknown identifier inside a node as the key (aliased or original) of an annotation that should be fetched for the item matched to this node. But since this is not always possible in an unambiguous way or there exists a name clash between an annotation key and another identifier available in the current environment, a $\langle \textit{qualifiedIdentifier} \rangle$ (snippet 2.4) can be used instead.

Grammar Snippet 2.4 ($\langle \textit{qualifiedIdentifier} \rangle$)

$\langle \textit{qualifiedIdentifier} \rangle ::= \langle \textit{identifier} \rangle ' :: ' \langle \textit{identifier} \rangle$

A qualified identifier consists of two identifiers that are joined by a double colon $' :: '$. The first (or left) identifier is expected to unambiguously resolve to a annotation layer available for the current context. The second (or right) identifier in turn acts as specification of an annotation key (or its alias) available from the previously defined annotation layer.


2.6.1.2 List Expression

Arrays (or more generally “lists”) in IQL can be defined by enclosing a sequence of expressions within curly brackets ($\{$ and $\}$), using the comma symbol ($,$) as separator. Optionally the desired type for elements in the list can be made explicit with a special type marker in front of the opening bracket as illustrated by snippet 2.5. If no type is provided, the overall element type will be determined by checking the types of all elements and choosing the least restrictive one. Note that for empty lists (no elements are given inside the curly brackets) it is mandatory to specify the element type!

Grammar Snippet 2.5 ($\langle \textit{listExpression} \rangle$)

$\langle \textit{listExpression} \rangle ::= (\langle \textit{type} \rangle ' [] ')? \{ (\langle \textit{expression} \rangle (',' \langle \textit{expression} \rangle)^*)? \}$

$\langle \textit{listExpression} \rangle$:



2.6.2 Path Expressions

For navigating hierarchically structured object graphs or namespaces, expressions can take the form of paths, consisting of a original expression, a dot as separator and finally an identifier that denotes the path element or “field” within the context of whatever the original expression returned.

Grammar Snippet 2.6 ($\langle path \rangle$)

$\langle path \rangle ::= \langle expression \rangle \text{ ' . ' } \langle identifier \rangle$

Examples:

```
someObject.someProperty
some.really.long.winded.path
```

Note that for a lot of native classes of the ICARUS2 framework, IQL provides convenient path-based alternatives to method invocations. For example in the context of navigating a structure, “someStructure.getParent(someItem)” can be replaced by “someItem.parent” as long as “someStructure” is unambiguous in the current context and already bound.

2.6.3 Method Invocation

Method invocations consist of an expression that points to the actual method (such as an identifier in the global namespace or a path expression) and round brackets for the invocation with an optional argument list:

Grammar Snippet 2.7 ($\langle method \rangle$)

$\langle method \rangle ::= \langle expression \rangle \text{ ' (' } \langle arguments \rangle \text{ ? ') ' }$

$\langle arguments \rangle ::= \langle expression \rangle \text{ (' , ' } \langle expression \rangle \text{) }^*$

$\langle method \rangle$:

$\rightarrow \langle expression \rangle - \text{ ' (' } \left[\begin{array}{c} \text{ ' , ' } \\ \langle expression \rangle \end{array} \right] \text{ ') ' } \rightarrow$

Examples:

```
myFunction()
myNamespace.someFunction(someArgument, anotherArgument)
min(123, 456, dynamicContent())
some().chained().methods()
```

2.6.4 List Access

Lists or arrays are accessed by an expression pointing to the list or array object itself and an index expression in square brackets indicating the position(s) of the desired element(s) within the array. Note that the index or indices expression must evaluate to values within 32bit signed integer space. Positive values indicate the position beginning from the start of the 0-based array, whereas negative values allow backwards referencing of elements with ‘-1’ pointing to the last array element and ‘-2’ to the second to last one. For multidimensional arrays several index statements can be chained or even combined in a single comma-separated list.

Grammar Snippet 2.8 ($\langle array \rangle$)

$$\langle array \rangle ::= \langle expression \rangle \text{ '[' } \langle indices \rangle \text{ ']' }$$

$$\langle indices \rangle ::= \langle expression \rangle (\text{' , ' } \langle expression \rangle)^*$$

$$\langle method \rangle :$$

$$\vdash \langle expression \rangle - \text{' [' } \overbrace{\langle expression \rangle}^{\text{' , ' }} \text{ ']' } \longrightarrow$$
Examples:

```
myArray[1]
myArray[-1]
myArray[-myArray.length] // same as myArray[0]
complexArray[1][2][3]
complexArray[-1][2][-3]
complexArray[1, 2][3]
complexArray[1, 2, 3]
```

Note that IQL provides convenient ways of using array access patterns to access list-like data structures and/or classes of the framework: Every ItemLookup implementation, such as Container or Structure that would traditionally access its content via “myContainer.getItemAt(someIndex)” can be used the same as any regular array with the expression “myContainer[someIndex]”.

2.6.5 Annotation Access

The ICARUS2 Corpus Modeling Framework (ICMF) models segmentation, structure and content of a corpus resource as different aspects. As such the information about any annotation attached to a given Item is stored apart from it and therefore is not easily accessible from the item alone. To simplify the usage of annotations within a query, IQL provides the following expression as syntactic sugar for accessing (multiple) annotations directly from an item:

Grammar Snippet 2.9 ($\langle annotation \rangle$)

$$\langle annotation \rangle ::= \langle expression \rangle \text{ '{ ' } \langle keys \rangle \text{ '}' }$$

$$\langle keys \rangle ::= \langle expression \rangle (\text{' , ' } \langle expression \rangle)^*$$

$$\langle method \rangle :$$

$$\vdash \langle expression \rangle - \text{' { ' } \overbrace{\langle expression \rangle}^{\text{' , ' }} \text{ '}' } \longrightarrow$$

The first expression must evaluate to an item reference and the annotation pointers inside curly brackets must evaluate to strings (if only a single expression is given, it can evaluate to a list or array and be expanded, cf. Section 2.6.13) that uniquely denote annotation layers in the current context of the query. Typically users will use string literals in double quotes to explicitly state the annotations to be accessed, but the IQL syntax allows for very flexible extraction statement. If the evaluation of those annotation pointers yields more than one string, the result will be an array-like object containing the resolved values for each of the annotation keys in the same order as those were specified.

Examples:

```
myItem{"pos"}
myItem{"form", "pos", "lemma"}

// extract values from multiple concurrent annotation layers
// and pick the first one present
firstSetValue(myItem{"parser1.head", "parser2.head"})
```

2.6.6 Type Cast

Expressions in IQL are automatically cast to matching types according to the actual consumer's needs (unless this feature gets deactivated via the corresponding switch, cf. Section 1.13.1). Explicit casts can be performed by preceding an expression with one of the type keywords listed above (2.1) in round brackets.

Examples:

```
(int) myValue
(int) 12345.678
(float) average(myVector)
(string) 123.456
```

2.6.7 Wrapping

Expression hierarchy and evaluation order follows the order the different types of expressions are listed here. To dictate another order, expressions can be wrapped into round brackets. This will cause the inner expression to be evaluated independent of potential hierarchical rules from the outside context.

Examples:

```
6 + 4 * 2    // multiplication evaluated first -> result 14
(6 + 4) * 2  // addition is evaluated first -> result 20
```

2.6.8 Unary Operation

IQL only allows four unary operators to be used directly in front of an expression, the exclamation mark '!' and the **NOT** keyword for Boolean negation, the minus sign '-' for negating numerical expressions and the '~' symbol of bitwise negation of integer numbers.

Examples:

```
!someBooleanFunction()
NOT someBooleanValue
-123
-myNumericalFunction()
~123
~myIntegerFunction()
```

2.6.9 Binary Operation

Binary operations between two expressions take the following simple form:

Grammar Snippet 2.10 ($\langle binary_op \rangle$)

$\langle binary_op \rangle ::= \langle expression \rangle \langle operator \rangle \langle expression \rangle$

Binary operators follow an explicit hierarchy, listed in Table 2.2 in the order of priority, from highest to lowest.

Operators	Explanation
* / %	multiplication, division and modulo
+ -	addition and subtraction
<< >> & ^	shift left, shift right, bitwise and, bitwise or, bitwise xor
< <= > >=	less, less or equal, greater, greater or equal
=~ !~ =# !#	string operators: matches (regex), matches not (regex), contains, contains not
== !=	equality, inequality
&& AND	logical conjunction
OR	logical disjunction

Table 2.2: Binary operators available in IQL and their hierarchical order.

2.6.9.1 Basic Numerical Operations

Basic numerical operations follow the standard mathematical rules for priorities. While the basic numerical types (**int**, **float**) can be arbitrarily mixed inside those expressions, the type used during the expression and as result will be determined by the least restrictive type of any operand involved.

2.6.9.2 Bit Operations

Bitwise operations (**&**, **|** and **^**) take integer expressions (or any other form of *bitset*) as inputs and generate a result of the corresponding type. If different types are used (e.g. **int** and **long**), one must be cast 2.6.6 to match the other. If value expansion 2.6.13 is active, any array-like data can also be used and will be subject to element-wise bit operations.

The two shift operations (**<<** and **>>**) take arbitrary integer types as left operand and an **int** value as right operand.

2.6.9.3 Ordered Comparisons

Comparisons are special binary operators that take two expressions of equal or compatible result type and produce a Boolean value. Note that their exact semantics are type specific, e.g. when comparing strings, the operation is performed lexicographically and may be subject to case conversions (1.13.1).

2.6.9.4 String Operations

To account for the ubiquity of textual annotations in corpora, IQL provides a set of dedicated string operators to perform substring matching (with the *contains* operator **=#** or its negated

form ‘!#’) and regular expression matching (via ‘=~’ and ‘!~’). Per default IQL uses the Java regex syntax, but for the future, additional switches (1.13.1) are planned to allow finer control over regex details.

Examples:

```
// find verbal forms
somePosAnnotation # "V"
// alternative to the set predicate with more flexibility
somePosAnnotation !~ "NN|NS"
```

2.6.9.5 Equality

Equality checks follow the same basic conditions as ordered comparisons (2.6.9.3), but with the following rules for comparable values “a” and “b”:

```
a == b iff !(a<b) && !(a>b)
a != b iff a<b || a>b
```

More generally, equality between expressions in IQL is based on content equality and therefore type specific. Note that trying to check two expressions of incompatible types (such as **int** and **string**) for equality will always evaluate to **false** and also emit a warning.

2.6.9.6 Logical Composition

All Boolean expressions can be combined via disjunction (either double pipes ‘||’ or the **OR** keyword) or conjunction (double ampersand ‘&&’ or the **AND** keyword), with conjunction having higher priority. While not strictly mandatory, evaluation of IQL expressions is recommended to employ optimized interpretation such that only the first operand is evaluated if possible. When the first operand of a disjunction evaluates to **true**, the entire expression is already determined, same for a conjunction’s first operand yielding **false**.

Examples:

```
a>1 && b<2
x==1 or x==3
```

Note that aforementioned optimization strategy also applies to chained boolean expressions, such as **a && b && c** and that the evaluation engine is free to reorder operands here as it sees fit. As such individual operands may never be evaluated at all and the overall correctness of a query expression must not rely on side-effects stemming from specific expressions being evaluated or not.

2.6.10 Assignment

IQL offers a special construct to assign (member) variables and at the same time verify the successful assignment with a Boolean result value. Snippet 2.11 illustrates the basic syntax for assignment operations inside a constraint, allowing the result of an arbitrary expression to be assigned to any previously defined member variable (or dynamically allocated variable).

$$\langle assignment \rangle ::= \langle expression \rangle \text{ 'AS' 'OPTIONAL' }? (\langle member \rangle / \langle variable \rangle)$$
$$\triangleright \langle expression \rangle - \text{'AS'}$$

OPTIONAL

$\langle member \rangle$

$\langle variable \rangle$

Examples:

2.6.11 Set Predicate

Set predicates can be directly negated (apart from wrapping 2.6.7 them and negating 2.6.8 the entire expression) with an exclamation mark ‘!’ or the keyword **NOT** in front of the **IN** keyword. If the input expression evaluates to an array-like object, the set predicate will expand its content and evaluate to **true** if at least *one* of its elements is found to be contained in the set. The set predicate can be universally quantified with a star ‘*’ or the **ALL** keyword in front to change the overall behavior such that the result will be **true** iff *all* of the elements are contained in the set (or none of them are, if the set predicate is directly negated).

$$\langle setPredicate \rangle ::= \langle expression \rangle \langle all \rangle? \langle not \rangle? \textbf{IN} \langle expression \rangle$$

$\langle all \rangle ::= \text{'ALL'} / \text{'*'}'$

$\langle not \rangle ::= \text{'NOT'} / \text{'!'}$

$\langle setPredicate \rangle:$

$\rightarrow \langle expression \rangle \underbrace{\text{'ALL'} / \text{'*'}}_{\text{'ALL'} / \text{'*'}} \underbrace{\text{'NOT'} / \text{'!'}}_{\text{'NOT'} / \text{'!'}} \text{'IN'} - \langle expression \rangle \rightarrow$

Examples:

```
someAnnotationValue IN {"NP","VP","-"}
someAnnotationValue NOT IN {"NN","DET"}
someAnnotationValue NOT IN {"NN","DET", STTS.getVerbTags()}
"John" IN getLegalNames()
fetchCharacterNamesInChapterOne() IN getOrcishNames()
```

2.6.12 Ternary Operation

A single ternary operation is supported in IQL, which is the popular if-then-else replacement with the following syntax:

Grammar Snippet 2.13 ($\langle ternary \rangle$)

$\langle ternary \rangle ::= \langle expression \rangle \text{'?' } \langle expression \rangle \text{' : ' } \langle expression \rangle$

The first expression must evaluate to a **boolean** value and determines which of the following two alternatives will be evaluated for the final value of the expression. Note that the second and third expressions must have compatible result types.

Examples:

```
x<2 ? "text for smaller value" : "some other text"
```

2.6.13 Value Expansion

IQL supports expansion of arrays, lists and array-like method return values for situations where an immediate consumer supports lists of values as input. Assuming the method “randomPoint()” returns an array of 3 integer values or a *array-like* data type (such as a 3D point) and another method “invertPoint(int, int, int)” takes 3 integer arguments, then the invocation of “invertPoint(randomPoint())” is legal and the array or object from the inner method call will be automatically expanded into the separate 3 values. This is especially handy when dealing with multidimensional arrays, as regular indexing would require manual extraction of method return values into variables to then be used in accessing the different array dimensions. With automatic expansion, a three-dimensional array could directly be accessed with aforementioned method via “array[randomPoint()]”.

2.7 Constraints

Simply put, constraints are expressions that evaluate to a Boolean result. Apart from native Boolean expressions (such as comparisons, Boolean literals or Boolean functions), IQL allows certain evaluations as syntactic sugar, listed in Table 2.3. Note that those conversions are only active if the respective switches to disable them (1.13.1) have not been set.

Type	Condition	Value
string	empty or null	false
int	0	false
float	0.0	false
any object	null	false

Table 2.3: Rules for converting arbitrary values or objects in a query to Boolean values.

Result Expansion Besides their obvious role as filters, constraints can also be used to expand the set of captured members in a match, subsequently usable for instance in advanced result processing.² The assignment operation (2.6.10) allows arbitrary expressions to be evaluated and having their results stored in either a general variable or, in case of a result type compatible to items, actual member variable (2.5).


2.8 Payload Structure

The Payload section in IQL consists of either the sole **ALL** keyword or a selection statement (2.8.4) with optional binding (2.8.3) definition and filter constraints (2.8.1) preceding it. If the **ALL** keyword is used, no constraints whatsoever can be defined and the engine is instructed to return the entire target corpus. In this case the only way of restricting results is by using the **iql:Result** section (1.17) of a query.

Grammar Snippet 2.14 (*<payload>*)

<payload> ::= 'ALL'
 / *<bindingsList>*? ('FILTER' 'BY' *<constraint>*)? 'FIND' *<modifier>*? *<selection>*
<modifier> ::= ('FIRST' / 'LAST' / 'ANY') (*<digits>* 'HITS')?

<payload with content>:



2.8.1 Filter Constraints

For complex (i.e. structural) queries, IQL offers a way of filtering the UoIs before they are processed by the matchers for sequence, tree or graph structures (cf. Sections 2.8.8 to 2.8.10). A dedicated **FILTER BY** section in the query payload preceding the actual structural constraints is available to define filtering rules that have to evaluate to **true** for a UoI to be considered for actual matching. Constraints within a filtering rules have only access to general properties of the UoIs, such as sentence length, tree height or similar information. The do **not** have access to bound member variables, apart from those defined for the top-level members of lanes (2.8.5)! Note that filter constraints are **not** compatible with flat constraints (2.8.6) as they both essentially fill the same function and flat constraints take precedence.

²Per default members defined as bindings (2.8.3) are captured when used as labels for nodes (2.8.7) in the query.

2.8.2 Match Modifiers

Per default, the search in IQL is expected to be exhaustive, i.e. the evaluation engine will attempt to find all of the instances in a target corpus that match the query constraints, potentially reporting individual UoIs (such as sentences) multiple times if they contain several occurrences. For instance, the sentence “The dog chased the rabbit down the hill.” will be treated as tree entries in the result if the query was only meant to find instances of the lemma “the”. This default behavior can be adjusted to only return each UoI no more than once by using one of the modifiers (**FIRST**, **LAST**, **ANY**) listed in snippet 2.14. The semantics of the **ANY** modifier are such that the engine may freely pick any one match within a UoI. Note however, that to support reproducible search results, repeated evaluations of the same query on a corpus are still required to yield the same instances here. The exact semantics of **FIRST** and **LAST** are depending on the type of structural constraints used in the payload, but generally are based on the natural order of items within the corpus (typically this is the flow of words in a text). The evaluation behavior for them is subsequently covered in Sections 2.8.8 to 2.8.10.

All three available modifiers can also take an optional numerical argument that defines the upper bound for the number of hits to be reported, followed by the **HITS** keyword. This *limit* must be a positive non-zero integer that follows the syntax of “pure integers” as described in Section 2.3.3. If a modifier is accompanied by a limit value, the ICQP will report a number of hits per UoI no greater than this limit.

2.8.3 Bindings

A binding is a collection of member references (2.5) that get declared to belong to a certain member type or part of the corpus. The **DISTINCT** keyword enforces that multiple bound member references in this binding do **not** match the same target. Depending on the local constraints used in the query, this might be redundant (e.g. when using the member references as identifiers for tree nodes who already are structurally distinct), but can still be used to make that fact explicit. Additionally the **EDGES** keyword signals that the bound members of a structure are actually edges. In this case using **DISTINCT** is redundant, as bound edges are implicitly assumed to be distinct when matching.


Grammar Snippet 2.15 (Bindings)

$\langle \text{bindingsList} \rangle ::= \text{WITH} \langle \text{binding} \rangle (\text{'AND'} \langle \text{binding} \rangle)^*$

$\langle \text{binding} \rangle ::= \langle \text{option} \rangle? \langle \text{member} \rangle (\text{' , ' } \langle \text{member} \rangle)^* \text{'FROM'} \langle \text{layer} \rangle$

$\langle \text{option} \rangle ::= \text{'DISTINCT'} / \text{'EDGES'}$

$\langle \text{bindingsList} \rangle:$



Raw binding definitions in the payload are parsed and stored in their JSON counterpart (`iq1:Binding`, 1.1) during query processing.


2.8.4 Selection Statement

Constraints are further divided into local constraints as part of node or edge definitions and global ones (with the **HAVING** keyword). Local constraints are obligatory and define the basic complexity of the query (flat, tree or graph). They also introduce certain limitations on what can be expressed or searched (e.g. a “flat” local constraints declaration will not provide implicit access to tree information). However, global constraints can introduce arbitrary constraints or relations and thereby increase the evaluation complexity, potentially without limits. Since there is no way for an evaluation engine to assess the complexity of user macros or extensions, extensive use of global constraints could in fact lead to extremely slow searches or even create situations where an evaluation will never terminate at all.


Grammar Snippet 2.16 ($\langle selectionStatement \rangle$)

$\langle selectionStatement \rangle ::= \langle constraint \rangle$
 $\quad | (\langle nodeStatement \rangle | \langle laneStatementsList \rangle) ('HAVING' \langle constraint \rangle)?$
 $\langle laneStatementsList \rangle ::= \langle laneStatement \rangle ('AND' \langle laneStatement \rangle)^*$
 $\langle laneStatement \rangle ::= 'LANE' \langle identifier \rangle ('AS' \langle member \rangle)? \langle nodeStatement \rangle$

$\langle selectionStatement \rangle$:



$\langle laneStatement \rangle$:



2.8.5 Lanes

Lane statements can be used to extract information from concurrent structures that exist for the UoI of the payload. Each lane statement is introduced by the **LANE** keyword and an identifier that matches the name or alias of a layer in the outer query definition (cf. snippet 2.16). Optionally the source layer of a lane can also be assigned a member variable (2.5) so that it can be explicitly referenced in the payload.³ During query processing raw lane statements will be parsed into `iql:Lane` objects (1.9).

Constraints are further divided into local constraints as part of node or edge definitions and global ones (with the **HAVING** keyword). Local constraints are obligatory and define the basic complexity of the query (flat, tree or graph). They also introduce certain limitations on what can be expressed or searched (e.g. a “flat” local constraints declaration will not provide implicit access to tree information). However, global constraints can introduce arbitrary constraints or relations and thereby increase the evaluation complexity, potentially without limits. Since there is no way for an evaluation engine to assess the complexity of user macros or extensions, extensive use of global constraints could in fact lead to extremely slow searches or even create situations where an evaluation will never terminate at all.

³This is particularly useful when using the global constraints to compare content of different lanes. Imagine for instance a query that searches for a certain syntactic construct **C** to be present in two concurrent parse trees **A** and **B**, but will only consider sentences where **C** is embedded deeper inside **A** compared to its embedding depth in **B**.

2.8.6 Flat Constraints

Flat constraints provide no extra helpers to declare structural properties of the query. They consist of arbitrary basic constraints Section 2.7 and disallow both global constraints (2.8.11) and filter constraints (2.8.1). Note that flat constraints rely on the availability of member references from the binding or lane sections in the query payload to have access to any content at all. In this regard they behave very similar to global constraints.

2.8.7 Structural Constraints

IQL provides several classes of structural constraints that each feature distinctive syntax features to express structures of increasing complexity. Those structures are sequences (2.8.8), trees (2.8.9) and graphs (2.8.10). They all get explained in more detail in their respective sections, but the syntactic basics for all of them will be defined here. To simplify the overall IQL grammar, a general syntax exists for the declaration of nodes (and edges). This general form honors the aspects specific to each of those structure types, but generally over-generates and only some of its features are actually applicable in concrete use cases. Snippet 2.17 shows the basic for defining structural constraints in IQL. More detailed illustrations of the various components ($\langle node \rangle$, $\langle quantifier \rangle$ and $\langle edge \rangle$) follow below.

Grammar Snippet 2.17 ($\langle structuralConstraint \rangle$)

$\langle structuralConstraint \rangle ::= \langle nodeStatement \rangle +$


$\langle nodeStatement \rangle ::= \langle quantifier \rangle? \{ ' \langle nodeStatement \rangle + ' \}$
 $/ \langle arrangement \rangle \langle nodeStatement \rangle +$
 $/ \langle node \rangle$
 $/ \langle element \rangle (' , ' \langle element \rangle)^*$
 $/ \langle nodeStatement \rangle \text{ 'or' } \langle nodeStatement \rangle$

$\langle arrangement \rangle ::= \text{ 'ORDERED' } / \text{ 'ADJACENT' }$


$\langle memberLabel \rangle ::= \langle member \rangle \text{ ':' }$

$\langle element \rangle ::= \langle node \rangle / \langle node \rangle \langle edge \rangle \langle node \rangle$

$\langle nodeStatement \rangle$:



$\langle element \rangle$:



2.8.7.1 Nodes

Grammar Snippet 2.18 ($\langle node \rangle$)

$$\langle memberLabel \rangle ::= \langle member \rangle \text{ ‘:’}$$
[illegible]

Grammar Snippet 2.19 (*quantifier*)

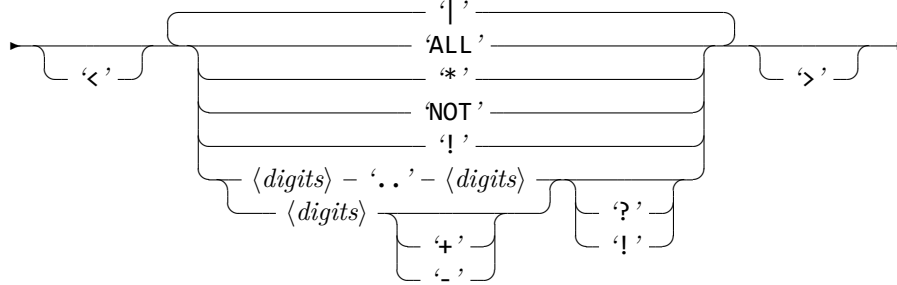
41

```

<simpleQuantifier> ::= ('ALL' / '*' )
/ ('NOT' / '!' )
/ <digits> ('+' / '-')? ('?' / '!' )?
/ <digits> '.' <digits> ('?' / '!' )?

```

<quantifier>:



A node's inner content can optionally have an initial **member label** (identifier with a colon ':' afterwards) to link this node to a previously defined member binding (2.8.3). Such as binding restricts the type of corpus member that the node can be matched against. It also provides a point of reference that subsequent constraints (e.g. in the global constraints section, cf. Section 2.8.11) can use to access information of the target the node has been matched against. Note that cross-referencing between nodes from within local constraints (such as in `[$x:][$y: pos!=$x.pos]`) is discouraged⁵ and global constraints should be used for this. This approach guarantees that by the time such cross-reference constraints (or "joins", to use database terminology) are evaluated all involved member variables will be assigned preliminary candidates.

Besides the local constraint described below, nodes can declare special **marker** expressions, separated from regular constraint or nested nodes by a single comma ','⁶. Marker can be regarded as glorified flags and functions that are used to specify a node's intended place in a sequence, hierarchy or tree path. They are logically separated from normal constraints, as they are used to construct an independent part of the ICQP state machine with the sole purpose of pruning the search space by limiting the actual target elements to be visited for individual nodes. A comprehensive introduction to the available markers of various types is given in Chapter 3.

Nodes can also optionally define **local constraints** that must evaluate to **true** for target item to be considered as result candidate. Local constraints have full access to properties and annotations of the target item currently being inspected and can take any form described in Section 2.7. Note that it is up to the evaluation engine how to optimize and potentially prune the evaluation process of constraint expressions. Correct evaluation of a conjunctive local constraints with external function calls `[$x: func1($x) && func2($x) && func3($x)]` must not rely on the premise of any particular function (`func1`, `func2` or `func3`) actually being

⁵The reason behind this is that the ICQP per default is not required to honor the order of nodes defined in a query or the linking relations between them when planning the state machine for evaluation. As such there is no guarantee that node `$x` will have already been matched against a valid target when the cross-reference constraint inside `$y` is evaluated. This would cause an error during evaluation time to occur, which in turn will abort the entire search.

⁶The IQL grammar declares this comma as mandatory, even if no local constraints or other contents succeeds the marker. This way they can be unambiguously distinguished from normal constraints, both by the ICQP and human readers.

called at all. The evaluation semantics of conjunctive Boolean concatenation allow an early determination of the final result as soon as the stable predicate of one of the inner terms evaluating to **false** is met. Therefore it is perfectly legal (and in parts expected from an efficient evaluation engine) to not evaluate the calls to **func2** and **func3** after **func1** has already caused the result to remain **false**.

If a $\langle node \rangle$ is used within a tree environment, it can also contain a **nested** $\langle nodeStatement \rangle$ declaration to define structural constraints on child nodes. The target item a node is matched against during query evaluation defines the structural context that is then in turn being used for matching the nested $\langle node \rangle$ instances. Note that this opens up the entire spectrum of node grouping, disjunctions and arrangement modifiers to define constraints on (sub-)trees.

2.8.7.2 Quantification

IQL knows a total of 12 quantifier types, which are grouped into six sets of basic quantification types as shown in Table 2.4. Those basic sets are *existential negation*, *universal quantification*, *exact*, *at-least*, *at-most* and *range*. Note that the *at-most* quantifier implies existential quantification, so it cannot quantify to less than one element. For modeling optional existence the *range* quantifier should be used with a lower bound of 0, as in **0..4**.

Some quantifiers differentiate between *greedy* (the default), *reluctant* and *possessive* mode. The differences of those three modes are explained below, assuming that the quantifier in question occurred on a node constraint n_i that is followed by an optional sequence of additional node constraints $n_{i+1}..n_N$. Note that the complete query will only succeed if n_i **and** all its successors successfully match! The primary effect caused by different modes is the size and number of matches reported by the evaluation engine.

greedy Match as many instances for n as possible, while still allowing $n_{i+1}..n_N$ to find appropriate candidates. This will find at least as many instances of n as the *reluctant* mode and at most as many as the *possessive* mode, but is guaranteed to leave room for $n_{i+1}..n_N$ when collecting hits.

reluctant Match as few instances for n as possible, considering the need of $n_{i+1}..n_N$ to also produce candidates. This will produce the shortest possible sequence of hits for n and is guaranteed to leave room for $n_{i+1}..n_N$ when collecting hits.

possessive Match as many instances for n as possible, completely ignoring $n_{i+1}..n_N$. This is guaranteed to find the longest sequence of hits for n but might cause the query to fail when subsequently there is no more room to match $n_{i+1}..n_N$.

An illustration of how the different quantification modes affect matching in an example sequence is shown in Fig. 2.1. Using two independent properties of the elements in a sequence (in this case color and shape) and an example query that uses them both a conflicting scenario is provided where individual elements can be matched by both query “nodes”. The Figs. 2.1b to 2.1d depict in simple terms how the different modes behave in such a conflicting situation.


The greedy mode (2.1b) will always produce the smallest amount of total matches of the entire query, as it is possible for adjacent matches to collapse. On the other hand, the reluctant mode (2.1c) will produce the maximum number of matches of an entire query, as it enforces the smallest individual match sizes. Finally Fig. 2.1d shows how the possessive mode can cause an entire query to fail by blindly consuming elements that otherwise could have been used by subsequent constraints to produce a successful match. If in the example however the node ‘8’ was square instead of round, the entire query could still have been successfully matched.

Greedy	Reluctant	Possessive	Description
not			Existential negation. Alternatively the exclamation mark ‘!’ can be used instead of the not keyword.
all			Universal quantification. Alternatively the asterisk ‘*’ can be used instead of the all keyword. Note that this quantifier can only be used directly on nodes and using it on a node that is not the only one within its context will cause an error.
X			Exact quantification of X instances.
X+	X+?	X+!	Open quantification of <i>at least</i> X instances.
X-	X-?	X-!	Bounded quantification of <i>at most</i> X instances.
X-Y	X-Y?	X-Y!	Bounded quantification of <i>at least</i> X , but <i>at most</i> Y instances.


Table 2.4: Complete list of element quantifiers supported by IQL, separated into the six fundamental groups. For quantifier types that do not distinguish between modes the *greedy* column shows the only way of formulating those. The effect of different modes is illustrated in Fig. 2.1.




(a) Target Sequence



(b) Hits for greedy mode: **red+** **box**.



(c) Hits for reluctant mode: **red+?** **box**. Since **red+?** allows **box** to match the first red square (5), the entire query can actually be matched twice.



(d) Partial hits for possessive mode: **red+!** **box**. Note that the query fails since **red+!** already consumes all the elements that could have been matched by **box**.


Figure 2.1: Graphical illustration of the different quantifier modes. The example query used is **red+** **box**, looking for first a sequence of red nodes and then a single square formed one.

2.8.7.3 Marker

Grammar Snippet 2.20 ($\langle \text{marker} \rangle$)

$\langle \text{marker} \rangle ::= \langle \text{identifier} \rangle ((' \langle \text{number} \rangle (' , ' \langle \text{number} \rangle) ^* ') ^?)$
 $/ (' \langle \text{marker} \rangle ')$
 $/ \langle \text{marker} \rangle (\text{'AND'} / \text{'\&\&'}) \langle \text{marker} \rangle$
 $/ \langle \text{marker} \rangle (\text{'OR'} / \text{'||'}) \langle \text{marker} \rangle$

$\langle \text{marker} \rangle$:


**2.8.7.4 Edges**

Edges are a structural element that is available exclusively to **graph** constraints (2.8.10). Each edge minimally consists of a source and target node and a type indicator⁷ to signal whether the edge is meant to be undirected, uni- or bidirectional and in the case of unidirectional edges which direction between source and target node it takes.

Grammar Snippet 2.21 ($\langle \text{edge} \rangle$)

$\langle \text{edge} \rangle ::= \text{'<--'} / \text{'-->'} / \text{'<->'} / \text{'---'}$
 $/ (\text{'<-'} / \text{'--'}) [' \langle \text{memberLabel} \rangle ? \langle \text{constraint} \rangle ? '] (\text{'--'} / \text{'>'})$

$\langle \text{edge} \rangle$:



Optionally, an edge can also be assigned its own member label and/or local constraints. In that case the type indicator gets split into two separate parts⁸ with a node-like part in the middle to host the edge's label and/or constraints. Note that the same guidelines for cross-referencing from within local constraints that were mentioned above in the section about nodes, also apply to local edge constraints: Correct evaluation of an edge's local constraints must not rely on the availability of cross-referenced external nodes (including the edge's own source and target nodes) or edges. As usual global constraints should be used to define constraints that link individual elements such as edges to other (external) parts of the payload. The section about global constraints (2.8.11) contains further hints on the optimization the ICQP implements in order to prevent unnecessary exploration of the search space.

⁷Type indicators for empty edges (i.e. edges without an inner constraint or label definition) always consist of three symbols, cf. snippet 2.21

⁸When split, each part of a type indicator uses two symbols, which are the respective 2/3 of the original three-symbol indicator.

The **quantification** of edges follows slightly different rules as compared to regular nodes. Every edge can have **up to one** explicit quantifier declaration attached to it, either on the source or target node. If no quantifier is present, the entire edge is by default existentially quantified, otherwise the following rules apply: The node not bearing any quantifier is existentially quantified (or “fixed”) and the quantifier on the other node acts as an indicator for the multiplicity of the edge itself. The following examples illustrate some use cases for quantification on edges:

```
// a node x linked to 3 nodes that match y
[$x:]--> <3>[$y:]
// 3 nodes matching x that are linked to the same node y
<3>[$x:]-->[$y:]
// node x with no direct link to a node that would match y
[$x:]-->![$y:]
```

Note that edge definitions cannot be chained, so the query “find x linked to y, linked to z” must be expressed in two separate edge definitions, one linking **\$x** to **\$y** and one linking **\$y** to **\$z**. This means that graph constraints (2.8.10) require some redundancy, but the linked section also explains some easy strategies to minimize overhead.

2.8.7.5 Grouping

Nodes (or elements) can be grouped together within curly brackets (‘{’ and ‘}’) as defined by the first *<nodeStatement>* rule in snippet 2.17. This is useful for either restricting the scope of modifiers or directives such as the **ADJACENT** keyword to only a selected few nodes or when expressing a disjunction. Optionally a node grouping can also be assigned quantifiers (2.8.7.2) similar to normal nodes (2.8.7.1) and as such can be used to express very powerful pattern such as repetition of sequences:

```
ADJACENT <3+>{ADJACENT [$a:][$b:]} [$c:]
```

The above query will match at least three repetitions of the pair a+b and the entire repeated sequence is followed by an adjacent node c.

Note that a group counts as an individual node statement inside the outer scope and as such is subject to order directives defined there. However, those directives are **not** automatically **inherited** to the inner collection of nodes in the group, allowing for expressions such as the following node sequence:

```
ADJACENT [$x:] {[$a:][$b:]} {[$c:][$d:]} [$y:]
```

This would read as “Find x immediately followed by a, later followed by b+c, later followed by d+y”. Note that the adjacency modifier does not apply to the inner sequences a+b and c+d, which are only subject to the implicit order of the sequence declaration.⁹ The concept of node grouping is especially important for the tree (2.8.9) and graph (2.8.10) constraints introduced below, as by default those do not impose an a priori order of nodes.

2.8.7.6 Node Sequence

Nodes usable for sequences (2.8.8) and trees (2.8.9) are defined in a simple sequence style (second *<nodeStatement>* rule in snippet 2.17). Instances of *<node>* in a sequence are defined

⁹An additional **ORDERED** in each of the groups would make that explicit, but is redundant.

one after another without special separator symbols. They may optionally be preceded by a *<arrangement>* directive to guide the matching progress. Currently there are only two directives available to specify the node arrangement (**ORDERED** and **ADJACENT**), but this might increase in the future, making node grouping a very important tool for defining complex structural compositions.

2.8.7.7 Element Sequence

Similar to node sequences, *<element>* instances can also be used in a list-style collection (third rule of *<nodeStatement>* in snippet 2.17), but with noticeable differences: Element sequences do use a separator symbol (a simple comma ‘,’) between *<element>* definitions. Since IQL does not use keywords to signal the structural type to be expected in a query payload¹⁰ this approach was necessary to easily detect the type of structure. It also hints at the second difference, that is, element sequences do not support arrangement modifiers (as *<element>* instances can be either nodes or edges, with the latter not being suitable for this kind of ordering) and as such can be more intuitively be understood as sets of *<element>* instances.

2.8.7.8 Structural Disjunction

As the forth option of *<nodeStatement>* in snippet 2.17, the disjunction of entire node statements provides a very powerful tool to express complex queries. Two node statements are linked with the **OR** keyword to mark them as “either or” cases. Since this opens up recursion, a complex disjunction may contain more than two node statements in total. Note that the order of node statements in a disjunction does not imply a specific priority the evaluation engine has to follow. For illustration, the second example below might match an instance of **[\$y:]** first:

```
[$x:] or [$y:]           // x or y
[$x:] or [$y:] or [$z]   // x or y or z
[$x:] or {[$y:][$z]}     // either x alone or a sequence y+z
[$x: {[$y:] or [$z]}]    // tree node x with either y or z as child
// complex nested disjunctions
[$x: [$y: [$z: [$a:] or [$b:]]] or [$z]]
```

2.8.8 Sequence Constraints

As the most basic form of structural constraints this type is used to match sequences of nodes to items in the target corpus. Multiple nodes in a sequence declaration are required to match to items in exactly the order they are defined in (but not necessarily adjacent to each other, use the **ADJACENT** directive in front of a node sequence for that).

Examples:

¹⁰An earlier draft made use of **TREE** and **GRAPH** keywords to distinguish those types from the basic node sequence, but in an effort to reduce the overall number of keywords (that users had to learn) this approach was dropped.


```

[] // empty node
[pos=="NN"] // node with local constraint
<2>[$x:] // node x exactly 2 times
[$x:]<2-5>[$y:] // nodes x & y with 2-5 nodes in between
[$x:] ![$y:] // node x without any node y following
ADJACENT [$x:][$y:] // node y directly following node x
[$x:] or [$y:] // disjunction: either x OR y
{[$x:][$y:]} or [$z:] // disjunction: either group x+y OR z

```

Sequence Matching Sequence constraints only provide a single dimension for *moving* the query sequence through the search space of the target corpus. Matching is performed greedily by default in order of node appearance in the query, following the direction specified by the corpus itself. That is, in a node sequence $[\$x:][\$y:]$ the node matched by $\$x$ will **always** be before the node matched by $\$y$ according to the original direction of items in the corpus. This also leads to a very simple and intuitive semantic for the **FIRST** and **LAST** modifier keywords: They stay true to their names and limit the returned matches to either the first or last ones, with respect to the current direction.

Empty nodes with quantifiers can be used as proxies to model distance constraints, as seen in above examples. Since the **ADJACENT** directive changes the behavior of an entire node sequence, some creativity can be necessary to achieve mixed cases, such as “*find an adjacent pair $a+b$ that is later followed by another adjacent pair $c+d$* ”. Possible (and simple) solutions for this query could be the following:

```

ADJACENT [$a:][$b:] <0+>[] [$c:][$d:]
{ADJACENT [$a:][$b:]} {ADJACENT [$c:][$d:]}

```

2.8.9 Tree Constraints

Located between mere sequences (2.8.8) and graphs (2.8.10) this type of structural constraints is meant to target tree structures, such as (but not limited to) syntax trees, coreference structures, discourse, etc. To simplify query syntax, IQL uses a similar approach as the original ICARUS project, which in turn took inspiration from PML-TQ: To signal parent-child relations, child nodes are nested within their respective parent, effectively making each node yet another scope for a sequence of child nodes. Contrary to bare sequence constraints (2.8.8) the order of (child) nodes to be matched in the corpus is **not** implicitly defined by the order of constraint nodes! Instead, the **ORDERED** or **ADJACENT** keywords need to be used explicitly to signal that a specific kind of order should be honored.¹¹ Apart from this little addendum tree constraints behave basically the same as nested sequence constraints: They can be individually quantified or existentially negated, as well as grouped and linked via the **OR** keyword to expression disjunctions.

Examples:

¹¹Also note that **ORDERED** and **ADJACENT** when used inside tree nodes are referring to the order of nodes as defined by the parent, **not** their overall position when projected onto the underlying foundation layer!

```

[[]]                // anonymous nesting of nodes
[$x: [$y:]]        // nesting of node y inside x
[$x:] [$y:]        // nesting of siblings x and y
[$x: [$y: [$z:]]]  // deep nesting chain
[$x: <2->[$y:]]     // at most 2 y nested inside x
[$x: ![$y:]]        // node x without any child matching y

// internal disjunction
[$x:] or [$y:] or {[$z: <4+>[]]}

```

Tree Matching As opposed to sequences (2.8.8), trees (and subsequently also graphs, cf. 2.8.10) offer an additional dimension of matching freedom over the bare iteration of items in a container or structure to be matched. This requires further specification of the matching process to guarantee consistent results and define basic expectations. Below explanations are primarily intended to specify the behavior in the presence of limiting modifiers (**FIRST**, **LAST**, **ANY**) in the payload, but are also of interest for the expected order of returned matches if no limiting modifiers are defined.




Figure 2.2: Example tree with highlighted hits for the simple tree query `[$x: [$y:]]`.

Consider a simple tree query `[$x: [$y:]]` where `$x` and `$y` are bound nodes with individual constraints. The nodes X_n or Y_n in the example tree (Fig. 2.2) then denote the n -th node that match the constraints of query nodes `$x` and `$y`, respectively. The nodes A_1 to A_6 are placeholders for nodes that match neither `$x` nor `$y`.

The tree contains 3 possible matches for the query, specifically the pairs $\{X_1, Y_1\}$, $\{X_1, Y_2\}$ and $\{X_2, Y_2\}$.
`[Root [A1] [A2 [X1 [A3] [Y1 [A4] [A5]] [Y2]]] [X2 [A6] [Y3]]]`

expand
on the
order of
matches

2.8.10 Graph Constraints

content, explain node and edge composition, mention referencing as a strategy to minimize edge declaration overhead, etc...

2.8.11 Global Constraints

Global constraints can be any basic constraint Section 2.7 and follow after the main section of structural constraints, indicated by the **HAVING** keyword (cf. snippet 2.16).

Evaluation Priorities If global constraints are present, the evaluation process changes to a two-stage strategy: Matchers for the associated structural constraints produce preliminary result candidates and the global constraints are then evaluated for each such candidate. This makes global constraints both very powerful as they have access to more information compared to regular (internal or local) constraints (e.g. they already *know* that all local constraints evaluated to **true** and the exact candidates produced for structural constraints) and also very critical when it comes to performance. It can be very tempting to construct queries such as the following one (bindings section omitted) that only matches when `$y` is the last child of `$x`:

```
FIND [$x: [$y:]] HAVING $x.indexOf($y) == $x.size-1
```

This will cause the structural matcher to potentially propose **all** children of **\$x** as candidates to be processed by the global constraints section. Subsequently, for a node of size N this will produce $N - 1$ candidates that are bound to fail the global constraint check. Chapter 3 lists several families of utility markers and functions that can be used to guide the evaluation engine into more efficient searches. With the use of those utility markers, above query looks like the following and will be vastly more efficient to evaluate:

```
FIND [$x: [$y: isLastChild,]]
```

Similarly global constraints are not the place to perform filtering on general properties of the current UoI, such as sentence length (use the **FILTER BY** expression for that, cf. Section 2.8.1).

Constraint Hoisting Per default, global constraints are second-class constraints, that are only consulted once the structural and local constraints in the other sections of a query payload have been evaluated. This provides them with a lot more decision power and available information, but also can lead to rather inefficient evaluation scenarios. Consider a query `FIND [$x: [$y: [$z]]] HAVING $x.pos!=$y.pos` that looks for nested nodes **\$x**, **\$y** and **\$z** where in addition to local constraints the nodes **\$x** and **\$y** are required to have different part-pf-speech tags. With the default two-stage evaluation process described above this would result in a serious overhead when evaluating a target tree: For every successful match of `[$x: [$y:]]` the engine would continue to look for children of **\$y** matching **\$z**, regardless of **\$x** and **\$y** satisfying the global constraint.

Through a optimization technique dubbed *constraint hoisting* the ICQP will try to work around this performance issue and attach global constraints to the nodes in the (tree) automaton that matches a certain part of the query. In the example above the matcher would evaluate the global constraint `$x.pos!=$y.pos` as soon as it successfully identified candidates for **\$x** and **\$y**, thereby reducing the overhead of searching for potential candidates for **\$z** that would ultimately fail due to **\$x** and **\$y** not satisfying the global constraint.

Note that only very specific global constraints can successfully be hoisted!

2.9 Result Processing

There be dragons...

(Content of the result section will be added as IQL evolves)

list
prop-
erties
required
for
hoist-
ing and
mention
hoisting
of con-
straint
frag-
ments,
since
IQL
splits
boolean
expres-
sions
into
frag-
ments

Chapter 3

Utility Markers & Functions

The following utility features are provided by the ICQP but are **not** part of the core specification. As such it is possible for engine extensions to override them, change their behavior or completely remove them if desired. They are listed here as per default they all are available and provide valuable improvements for performance and usability.

3.1 Position Markers

Every node (2.8.7) in a query has an implicit Container or Structure context that it is hosted or contained in.¹ IQL offers a variety of predefined helper functions to exploit this fact and to optimize queries. Table 3.1 lists the functions currently available, their arguments and matching conditions. Note that legal arguments for a node that occurs in a container of size N reside in the closed integer interval $[1..N]$ for regular index values and $[-N..-1]$ for reverse indices (see below).² The syntax for providing arguments follows the normal rules for method invocations in IQL 2.6.3, e.g. **isOutside(4, 12)**. Arguments can be provided as literals, which is the preferred way, or as any other kind of expression that evaluates to the necessary type (or can be converted, depending on the query configuration). It is not possible to access properties from the node itself for which the marker is defined, as position markers are evaluated **before** the corresponding nodes are matched and subsequently any such attempts will result in an error during query evaluation. However, globally accessible information or any data provided by surrounding containers or structures can be used to specify marker arguments.³

3.1.1 Negation

Position markers cannot be negated when used in a constraint expression.⁴ For every function/marker in Table 3.1 there exists a complement that can be used to easily express the negated condition. For instance **isBefore(n)** can be negated into **isAfter(n-1)**, but the engine will not perform this conversion and it up to the user to properly provide the intended

¹The ICMF specifies that each item can only be hosted (or more accurately *owned*) by a single container or derived object, but be contained within an arbitrary number of additional containers or structures.

²This is in contrast to the 0-based access of array or list elements in IQL.

³For example, a node could be defined to be in first position for short sentences and in last position for long ones.

⁴Depending on the query configuration doing so would either fail the evaluation process with an error or silently ignore the issue, potentially producing invalid results.

Label	Arguments	Matching Condition
isFirst	-	Node is the first in its context ($index = 1$).
isLast	-	Node is the last in its context ($index = N$).
isAfter	int	Node occurs after designated position ($index > arg$). Legal values are: $1 \leq arg < N$ (or $-N \leq arg < -1$ for reverse indices).
isBefore	int	Node occurs before designated position ($index < arg$). Legal values are: $1 < arg \leq N$ (or $-N < arg \leq -1$ for reverse indices).
isInside	int, int	Node occurs inside designated interval ($arg_1 \leq index \leq arg_2$). Legal values are: $1 \leq arg_1 < arg_2 \leq N$ (or $-N \leq arg_1 < arg_2 \leq -1$ for reverse indices).
isOutside	int, int	Node occurs outside designated interval ($index < arg_1$ and $index > arg_2$). Legal values are: $1 \leq arg_1 < arg_2 \leq N$ (or $-N \leq arg_1 < arg_2 \leq -1$ for reverse indices).
isAt	int	Node occurs at specific position ($index = arg$). Legal values are: $1 \leq arg \leq N$ (or $-N \leq arg \leq -1$ for reverse indices).
isNotAt	int	Node occurs at any position except the designated one ($index \neq arg$). Legal values are: $1 \leq arg \leq N$ (or $-N \leq arg \leq -1$ for reverse indices).

Table 3.1: Position markers provided by the default evaluation engine for IQL. The conditions in the last column assume a container of size N as context. If a marker function takes arguments, they appear in the conditions as arg for the sole argument or arg_i for the i -th argument list, with arg_1 being the first argument.

negated form. Trying to negate a position marker inside structural constraints will result in an error during the preprocessing of the query payload.

3.1.2 Reverse Indices

Any position marker that takes at least one argument also supports negative values that are treated as reverse indices, starting from the end of the surrounding container’s size. This follows the same semantics as the index scheme for arrays or lists (2.6.4). Note that certain position markers with interval arguments (such as `inBetween(arg_1 , arg_2)`) still require actual interval boundaries to be in the proper order (generally $arg_1 \leq arg_2$).

3.1.3 Relative Indices

Per default all the position markers that take argument expect them to be integer values for explicit designation of the desired index or index range. This means that any other type of argument will be automatically converted to an integer value if the query configuration permits it. However, if the associated switch (1.13.1) is active⁵, position markers can use relative

⁵This switch is off by default as it interferes with the automatic type conversion of arguments. If it is active, arguments must explicitly be cast if not already an integer or floating-point number. If the switch is not active, any floating-point argument will be cast to an integer (typically 0 if any of the arguments are intended to be reverse indices).

arguments in the form of floating-point numbers. Relative indices are expected to be in the open interval (0..1) (or (-1..0) for negative reverse indices) and are treated as percentages of the total size of the surrounding container. The following examples illustrate this approach (assume the markers are used inside a token node that is directly hosted within a sentence):

```
// the token can only occur in the first half of the sentence
isBefore(0.5)
```

```
// the token must occur in the closed interval [20%, 80%]
isBetween(0.2, 0.8)
```

The actual index values to be used for determining the legal position of the node are calculated by multiplying the context's size attribute (e.g. the length of the surrounding sentence) with the marker's relative argument and then rounding up the result to the nearest integer, after which a conversion into 0-based index space is applied. It is easily possible to define the same position marker with relative indices in both normal and reverse form. For instance, `isBefore(0.8)` and `isBefore(-0.2)` both describe the same 80% interval. As a general rule, relative indices should only be used with regular (positive) arguments.

3.1.4 Marker Stacking

Position markers can freely be mixed in disjunctive or conjunctive expressions.

```
[isLast,]
[isBefore(10), pos!="NN"]
[isBefore(10) || isAfter(24), pos!="NN"]
[isNotAt(5) && (isBefore(10) || isAfter(24)), pos!="NN"]
```

The reasons for this stems from the underlying interval arithmetic. Each positional marker effectively describes one or more non-empty intervals of possible index values for the node based on the overall size of the target container. Conjunction of markers results in the intersection of their intervals, possibly creating an empty set, which will render the node constraint impossible to satisfy. Disjunction of markers on the other hand creates a union of their intervals, either joining them if they overlap or creating a set of disjoint intervals that denotes a discontinuous collection of legal values. The final result is a fixed⁶ collection of index values that can be iterated to check for potential result candidates. In the absence of any kind of positional marker this default index set to be traversed is the complete interval [1..N] for a host container of size *N*.

rework,
since we
dropped
the
DNF
require-
ment

3.1.5 Markers & Quantification

Markers can be freely combined with quantifiers on the same node, with the following rules:

- Markers apply to **all** instances of a node, so they effectively take priority over quantifiers that allow multiple instances of a node to match.
- A node that bears at least one marker is effectively existentially quantified and so it is **not** possible to combine markers and the *existential negation* quantifier.

⁶Fixed in the sense that it only depends on the overall size of the surrounding container.

- In similar fashion, universal quantification is not compatible with markers of any kind, even if they would allow the entire target sequence as legal index interval.

To illustrate the interaction of quantifiers and markers, consider the example of a node `<3..5>[$x: isInside(1,0.25),]`. This query requires at least three, but no more than five instances of `$x` to be matched in the first 25% of a target sequence. Evaluated against a sequence of size 10, this query will automatically fail, as the legal interval produced by the marker only provides access to the first two elements of the sequence⁷. A sequence of size 12 is the smallest target that even has a chance of being matched. Note however that until the target size for this example reaches 20, the total number of hits for `$x` will never satisfy the upper bound for the range quantifier, as the marker’s index interval won’t support enough elements.

3.1.6 Markers in Global Constraints

Global constraints (2.8.11) have access to an extended version of the “local” position markers that acts as a general predicate for items (and additional index arguments, depending on the marker). So for example `isLast($x)` is the method-equivalent to the local marker `[$x: isLast,]`. Note that there is no performance benefit by using those methods in global constraints, as actual position markers are only recognized within structural constraints. The version for global constraints merely serves as syntactic sugar. However, since those methods are not subject to the same limitations as mentioned in the section about stacking above, they can be used freely inside local constraints.

3.2 Tree Markers

While the basic positional markers described in Section 3.1 limit a node’s overall position in the surrounding container, tree markers work on the hierarchical properties of tree structures. As such they define legal positions of a node within its parent’s list of children. All tree markers are subject to the same rules as position markers regarding negation, relative and reverse indices, as well as stacking. Note however, that the default order for traversing child nodes in a tree does not necessarily⁸ reflect the actual order of the child nodes position when projected onto the underlying foundation layer. It can therefore make sense to define additional precedence constraints via global constraints if needed.

3.2.1 Hierarchy Markers

Apart from the horizontal domain described above, tree markers also cover the hierarchical aspect of tree structures. A series of additional markers is available to specify a node’s vertical location in the tree, as shown in Table 3.3. These markers inherit all the rules of general tree markers, except that they cannot be used with relative indices.

All the additional hierarchical markers in Table 3.3 operate purely on the vertical axis within a tree structure (with the exception of the two “descendants” markers that use the fringes of the subtree spanned by a node’s parent to place candidate nodes in the target tree). The initial lot of them (`isRoot`, `isNoRoot`, `isLeaf`, `isNoLeaf` and `isIntermediate`) are

⁷Markers round down to the nearest integer size for relative indices, see Section 3.1.3.

⁸The ICMF does however strongly encourage tree structures to order children according to their begin-offsets if it does not interfere with their intended semantics.

Label	Arguments	Matching Condition
isFirstChild	-	Node is the first child of its parent ($childIndex = 1$).
isLastChild	-	Node is the last child of its parent ($childIndex = N - 1$).
isChildAfter	int	Node occurs as child after designated position ($childIndex > arg$). Legal values are: $1 \leq arg < N$.
isChildBefore	int	Node occurs as child before designated position ($childIndex < arg$). Legal values are: $1 < arg \leq N$.
isChildInside	int, int	Node occurs as child inside designated interval ($arg_1 \leq childIndex \leq arg_2$). Legal values are: $1 \leq arg_1 < arg_2 \leq N$.
isChildOutside	int, int	Node occurs as child outside designated interval ($childIndex < arg_1$ and $childIndex > arg_2$). Legal values are: $1 \leq arg_1 < arg_2 \leq N$.
isChildAt	int	Node occurs as child at specific position ($childIndex = arg$). Legal values are: $1 \leq arg \leq N$.
isChildNotAt	int	Node occurs as child at any position except the designated one ($childIndex \neq arg$). Legal values are: $1 \leq arg \leq N$.
isLeftChild	-	Node's covered region on the foundation layer is outside and to the left of that covered by its parent ($child.right < parent.left$). Legal values are: $1 \leq arg \leq N$.
isRightChild	-	Node's covered region on the foundation layer is outside and to the right of that covered by its parent ($child.left > parent.right$). Legal values are: $1 \leq arg \leq N$.

Table 3.2: Positional markers for trees provided by the default evaluation engine for IQL. The conditions in the last column assume a parent node with N children as context and $childIndex$ being the desired position of the node within its parent's list of children. If a marker function takes arguments, they appear in the conditions as arg for the sole argument or arg_i for the i -th argument list, with arg_1 being the first argument.

based on the presence or absence of child nodes and/or a parent. The remaining markers use the concept of *generations*⁹ in the tree. Given a tree structure T and the root node R , the set of nodes in T are partitioned into disjoint sets of generations G_0 to G_n where n is the height of T and G_i is the set of nodes belonging to the i -th generation, that is all the nodes with a depth of i (i.e. path distance i to R). Note that R is the only member in G_0 and G_n can only contain leaf nodes. Generations in structural IQL queries are always calculated based on the immediate parent of the node declaration that contains the marker(s). If a constraint related to the nesting distance between transitively nested nodes or other non-immediate relations is desired, global constraints (2.8.11) should be used, possibly employing dedicated tree functions (3.5).

If no generation markers are used within a node constraint and the node is nested, a direct parent-child relation is assumed, which is equivalent to **isGeneration(1)**. Nodes that represent the root of a tree query can also use generation-based markers to indicate where in a

⁹Also often called “levels” in tree structures.

Label	Arguments	Matching Condition
isRoot	-	Node is a designated root in the structure.
isNoRoot	-	Node has a parent.
isLeaf	-	Node is a leaf, i.e. it has no children.
isNoLeaf	-	Node is not a leaf, i.e. it has at least one child.
isIntermediate	-	Node is neither a designated root nor a leaf.
isGeneration	int	Node is a member of the <i>arg</i> -th generation of the parent node. Note that only values of $arg \geq 2$ make any sense, as per default parent-child nesting in tree constraints implies membership in the 1 st generation.
isNotGeneration	int	Node is nested arbitrarily deep in its parent, but not in the designated generation.
isGenerationAfter	int	Node is nested at least $arg + 1$ generations deep in its parent.
isGenerationBefore	int	Node is nested no more than $arg - 1$ generations deep in its parent.
isAnyGeneration	-	Node is nested arbitrarily deep in its parent. This is the equivalent of a full transitive closure over the tree dominance relation between parent and child nodes.
isLeftmostt	-	Node is located along the path of leftmost descendants (according to the actual position on the underlying foundation layer, not necessarily the order of child nodes along each step). This marker is generally combined with one of the generation-based markers to provide an indicator on the vertical location of the desired descendant.
isRightmost	-	Represents the symmetric opposite of above isLeftmost , marker, following the path of rightmost descendants.

Table 3.3: Additional markers for hierarchical properties in tree structures. Note that “parent” in the matching conditions refers to the node in a target tree that has been matched by the original node’s parent.

prospective target tree the associated item should be located. Regular tree hierarchy markers can also be combined with generation-based markers to further specify the location within a tree. For instance, [**\$x**: [**isGenerationAfter(2) && isNoLeaf,**]] will match any node that is nested at least 3 steps deep within **\$x** but ignores leaves. The **isRoot** marker however cannot be combined with any other hierarchy-related markers, as only designated roots in the target structure can match it and the only sensible combinations would be with normal (horizontal) position markers (3.1). If the **isRoot** marker is being used within a nested node, an error will be issued as the query will be impossible to satisfy.¹⁰

¹⁰The ICQP does this for any kinds of obviously erroneous queries when the planning phase results in unsolvable issues.




Figure 3.1: Example tree with highlighted nodes on the fringes for the **isLeftmost** (green, nodes 1, 4 and 11) and **isRightmost** (red, nodes 3 and 10) markers.

3.2.2 Path Markers

While hierarchical and regular tree markers (3.2, 3.2.1) operate on strictly one of either the horizontal or vertical dimension within a tree, path markers effectively combine horizontal and vertical navigation into a single marker. As such they describe a path starting from a particular node downwards rather than a specific position within that node’s list of immediate children or a vertical range within the tree.

Label	Arguments	Matching Condition
isLeftmost	-	Node is located along the path of leftmost descendants (according to the actual position on the underlying foundation layer, not necessarily the order of child nodes along each step).
isRightmost	-	Represents the symmetric opposite of above isLeftmost , marker, following the path of rightmost descendants.

Table 3.4: Additional markers for path properties in tree structures.

Fringe Markers The two **fringe** markers **isLeftmost** and **isRightmost** signal that node candidates must be located at the left (or right) fringe of a target tree as illustrated by Fig. 3.1. Note that the *leftmost* and *rightmost* properties do **not** refer to the position of children within their parent’s list of child nodes. Instead the actual position on the underlying foundation layer is taken as basis for determining which nodes belong to the respective fringes. More specifically, an item’s begin index is responsible for its position in the lift of *leftmost* candidates, and the end index is taken as the decisive property when it comes to the *rightmost* fringe. Since the ICMF does not require structures to order outgoing edges for a parent node according to the positional indices of its children, using those fringe markers can incur a certain computational overhead.

Quantification With regular tree markers (3.2) quantifiers operate along the horizontal axis (either globally or within the context of a node’s list of children). Path markers however completely change this behavior in a way that quantification happens along the path defined

by the marker. This means that quantification effectively changes to be vertical with the path marker acting as guide for picking the next path element from a horizontal pool of candidates.

Marker Stacking Path markers are typically combined with one or more generation-based markers to further specify a node’s position in the tree. If no additional hierarchical marker is provided only the immediate children of a node are taken into account. Otherwise the path marker will guide the selection of child nodes along the path to the desired vertical position. Note however, that the presence of quantifiers greatly restricts what kind of marker stacking is actually legal in a particular case. If for instance a quantifier requires more than one instance of the node to be matched along the path, but a generation-based marker limits the occurrence to be on only a single level, the query becomes unsolvable and the ICQP will issue an error.

Examples: ¹¹

```
[isLeftmost && isLeaf,]           // node 11
[isRightmost && isGeneration(-2),] // node 3
// go down 2 to 4 nodes alongside the left fringe
<2-4>[isLeftmost,]
// impossible query
<2>[isLeftmost && isGeneration(3),]
```

3.2.3 Evaluation Performance

Tree-related methods for fetching hierarchical information (such as height and depth of a node) are optional in ICMF. As a result the ICQP might have to compute the data required for evaluating a query at runtime, generating some additional overhead from tree traversals. This is especially critical for hierarchical markers that use reverse indices, as knowledge about the overall tree height is required to effectively translate relative indices into absolute intervals of legal (generation) values.

3.3 Sequence Functions

finish

correct
descrip-
tion/-
caption
in table

¹¹Node references based on the example tree in Fig. 3.1.

Label	Arguments	Result	Description
isAdjacent	node[]	Boolean	Check whether the given nodes form a continuous interval on the underlying foundation layer.
isOrdered	node[]	Boolean	Check whether the given nodes form an ordered sequence, as indicated by their begin indices, allowing for nodes to overlap.
isOrderedDisjoint	node[]	Boolean	Check whether the given nodes form a disjoint ordered sequence, as indicated by their begin indices.

Table 3.5: Utility functions for advanced sequence matching. Primarily these functions are intended to provide equivalents for the position markers in Section 3.1 usable in global constraints and batch methods for checking the horizontal arrangement of nodes.

3.4 Spatial Functions

To provide fine control over horizontal positioning of nodes, a set of utility functions is provided that models specific spatial relations regarding (relative) positioning, overlap or inclusion of a pair of nodes.

Label	Matching Condition
isLeftOf	Node n_1 located completely outside and to the left of n_2 ($n_1.right < n_2.left$).
isRightOf	Node n_1 located completely outside and to the right of n_2 ($n_1.left > n_2.right$).
overlaps	Nodes n_1 and n_2 overlap in some way ($n_1.left \leq n_2.right \wedge n_2.left \leq n_1.right$).
overlapsNot	Nodes n_1 and n_2 do not overlap ($n_1.left > n_2.right \vee n_2.left > n_1.right$). This function is a negation of overlaps that mainly exists as syntactic sugar.
overlapsLeft	Node n_1 overlaps n_2 on the left side ($n_1.left \leq n_2.left \wedge n_1.right \geq n_2.left$).
overlapsRight	Node n_1 overlaps n_2 on the right side ($n_1.right \geq n_2.right \wedge n_1.left \leq n_2.right$).
surrounds	Node n_1 fully surrounds n_2 , i.e. it overlaps n_2 on both sides ($n_1.left \leq n_2.left \wedge n_1.right \geq n_2.right$).
fits	Node n_1 exactly fits n_2 , i.e. its boundaries match on both sides ($n_1.left = n_2.left \wedge n_1.right = n_2.right$).
alignsLeft	Node n_1 and n_2 share the same left boundary ($n_1.left = n_2.left$).
alignsRight	Node n_1 and n_2 share the same right boundary ($n_1.right = n_2.right$).

Table 3.6: Spatial functions to model relative positioning between two nodes. All functions listed take exactly two node arguments, n_1 and n_2 . For the matching condition column, $n.left$ stands for the begin index (the left boundary of the span covered by n when projected onto the common foundation layer) and $n.right$ for the end index with $n.left \leq n.right$.

3.5 Tree Functions

Label	Arguments	Result
ancestor	node[]	Returns the first common ancestor of the given nodes.
parentAt	node, int	Returns the arg_2 -th parent of arg_1 . The immediate parent of arg_1 is reached with $arg_2 = 1$, its grandparent with $arg_2 = 2$ and so on ...Note that a node's immediate parent can (in the presence of a single unambiguous structure) be reached with the parent shortcut field.

Table 3.7: Utility functions for advanced tree matching. Primarily these functions are intended to complement the top-down tree matching strategy employed by the ICQP with bottom-up evaluation and navigation options.

The nested style of defining tree queries in IQL in top-down manner and the associated matcher in the ICQP also producing top-down automata for evaluation cause some obvious shortcomings: For problems in bottom-up style, for instance “Find the first common ancestor of adjacent phrases x and y ” it can be more intuitive (and usually more efficient wrt evaluation) to define the query bottom-up. Using the tree functions listed in Table 3.7 a simple query for this question could look like the following:

```
FIND ADJACENT [$x:][$y:] HAVING ancestor($x,$y) AS $a}
```

This will first try to find two adjacent instances for $\$x$ and $\$y$ from the set of nodes in the target tree and then use the `ancestor($x,$y)` method to fetch their first common ancestor node and assign it to $\$a$ for further use. The same can be achieved in a top-down fashion with a slightly more verbose query:

```
FIND [$a: childCount>=2
    [$x: isAnyGeneration,]
    [$y: isAnyGeneration,]]
HAVING isAdjacent($x, $y)
```

Besides the differences in appearance those two approaches also differ greatly in their respective evaluation complexity. While the bottom-up query can be efficiently solved in $\mathcal{O}(n + h)$ time for trees with n nodes and height h ¹², the top-down query results in a vastly more expensive search. Since node $\$a$ does not provide any strong filter mechanism besides needing at least 2 child nodes, the matcher will have to move the search tree through the target tree and repeatedly check any possible combination of descendant nodes matching $\$x$ and $\$y$ to be adjacent.

3.6 Graph Functions

There be dragons...

(Content of the this section will be added as IQL evolves)

¹²Construction of a reverse lookup between indices and the nodes beginning/ending there and subsequent production of adjacent pairs for $\$x$ and $\$y$ requires $\mathcal{O}(n)$ time. Without a dedicated utility data structure, determining the first common ancestor of $\$x$ and $\$y$ is linear in the maximum depth of the two nodes.

Appendices

Appendix A

Extended Grammar Diagrams

Grammar Snippet A.1

$\langle expression \rangle$:

